NABIN KATWAL

MATHEMATICS
FOR
DATA SCIENCE

Contents

Nabin Katwal

Nabin Katwal

Nabin Katwal

Nabin Katwal   7

7 CLUSTERING   163

8 DIMENSIONALITY REDUCTION 189

Nabin Katwal

Nabin Katwal   11

Chapter 1

LINEAR ALGEBRA

1.1.   Vector Addition

$$\mathbf{u} + \mathbf{v} = \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix} + \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} = \begin{bmatrix} u_1 + v_1 \\ u_2 + v_2 \\ \vdots \\ u_n + v_n \end{bmatrix}$$

Explanation: Vector addition combines two vectors component-wise. It is commonly used in machine learning for gradient updates or geometric vector operations.

Example: If $\mathbf{u} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$ and $\mathbf{v} = \begin{bmatrix} 3 \\ 4 \end{bmatrix}$, then $\mathbf{u} + \mathbf{v} = \begin{bmatrix} 4 \\ 6 \end{bmatrix}$.

Implementation:

```
import numpy as np
u = np.array([1, 2])
v = np.array([3, 4])
result = u + v
```

1.2. Scalar Multiplication of a Vector

$$\alpha \mathbf{v} = \alpha \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} = \begin{bmatrix} \alpha v_1 \\ \alpha v_2 \\ \vdots \\ \alpha v_n \end{bmatrix}$$

Explanation: Scalar multiplication scales each component of a vector by the same scalar. It is used in scaling gradients or controlling vector magnitudes.

Example: If $\alpha = 3$ and $\mathbf{v} = \begin{bmatrix} 2 \\ -1 \end{bmatrix}$, then $\alpha \mathbf{v} = \begin{bmatrix} 6 \\ -3 \end{bmatrix}$.

Implementation:

```
import numpy as np
alpha = 3
v = np.array([2, -1])
result = alpha * v
```

1.3. Dot Product

$$\mathbf{u} \cdot \mathbf{v} = \sum_{i=1}^{n} u_i v_i = u_1 v_1 + u_2 v_2 + \cdots + u_n v_n$$

Explanation: The dot product calculates a scalar representing the magnitude of projection of one vector onto another. It is widely used in ML for similarity measures or linear operations.

Example: If $\mathbf{u} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$ and $\mathbf{v} = \begin{bmatrix} 3 \\ 4 \end{bmatrix}$, then $\mathbf{u} \cdot \mathbf{v} = 1 \cdot 3 + 2 \cdot 4 = 11$.

Implementation:

```
import numpy as np
u = np.array([1, 2])
v = np.array([3, 4])
result = np.dot(u, v)
```

Nabin Katwal

1.4. Cross Product (3D)

$$\mathbf{u} \times \mathbf{v} = \begin{vmatrix} \mathbf{i} & \mathbf{j} & \mathbf{k} \\ u_1 & u_2 & u_3 \\ v_1 & v_2 & v_3 \end{vmatrix}$$

Explanation: The cross product generates a vector perpendicular to two input vectors in 3D space. It is commonly used in physics and computer graphics.

Example: If $\mathbf{u} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$ and $\mathbf{v} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$, then $\mathbf{u} \times \mathbf{v} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$.

Implementation:

```
import numpy as np
u = np.array([1, 0, 0])
v = np.array([0, 1, 0])
result = np.cross(u, v)
```

1.5. Norm of a Vector (Euclidean)

$$\|\mathbf{v}\| = \sqrt{\sum_{i=1}^{n} v_i^2} = \sqrt{v_1^2 + v_2^2 + \cdots + v_n^2}$$

Explanation: The Euclidean norm measures the magnitude (length) of a vector. It is useful in optimization and distance computations in ML.

Example: If $\mathbf{v} = \begin{bmatrix} 3 \\ 4 \end{bmatrix}$, then $\|\mathbf{v}\| = \sqrt{3^2 + 4^2} = 5$.

Implementation:

```
import numpy as np
v = np.array([3, 4])
result = np.linalg.norm(v)
```

1.6. Orthogonality Condition

$$\mathbf{u} \cdot \mathbf{v} = 0$$

Explanation: Two vectors are orthogonal if their dot product is zero. This condition is critical in linear algebra and ML for understanding independence and basis construction.

Example: If $\mathbf{u} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$ and $\mathbf{v} = \begin{bmatrix} -2 \\ 1 \end{bmatrix}$, then $\mathbf{u} \cdot \mathbf{v} = 1 \cdot -2 + 2 \cdot 1 = 0$, confirming orthogonality.

Implementation:

```
import numpy as np
u = np.array([1, 2])
v = np.array([-2, 1])
result = np.dot(u, v)
is_orthogonal = result == 0
```

1.7. Matrix Addition

$$\mathbf{A} + \mathbf{B} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} + \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{bmatrix} = \begin{bmatrix} a_{11} + b_{11} & a_{12} + b_{12} \\ a_{21} + b_{21} & a_{22} + b_{22} \end{bmatrix}$$

Explanation: Matrix addition combines two matrices element-wise. It is used in ML for updating weights and biases or aggregating data.

Example: If $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $\mathbf{B} = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$, then $\mathbf{A} + \mathbf{B} = \begin{bmatrix} 6 & 8 \\ 10 & 12 \end{bmatrix}$.

Implementation:

```python
import numpy as np
A = np.array([[1, 2], [3, 4]])
B = np.array([[5, 6], [7, 8]])
result = A + B
```

1.8.   Matrix Scalar Multiplication

$$\alpha \mathbf{A} = \alpha \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} = \begin{bmatrix} \alpha a_{11} & \alpha a_{12} \\ \alpha a_{21} & \alpha a_{22} \end{bmatrix}$$

Explanation: Scaling a matrix by a scalar is useful in ML for adjusting learning rates or normalization.

Example: If $\alpha = 2$ and $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$, then $\alpha \mathbf{A} = \begin{bmatrix} 2 & 4 \\ 6 & 8 \end{bmatrix}$.

Implementation:

```
import numpy as np
alpha = 2
A = np.array([[1, 2], [3, 4]])
result = alpha * A
```

Nabin Katwal

1.9.   Matrix-Vector Multiplication

$$\mathbf{Ax} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} a_{11}x_1 + a_{12}x_2 \\ a_{21}x_1 + a_{22}x_2 \end{bmatrix}$$

Explanation: Matrix-vector multiplication transforms a vector using a linear transformation defined by the matrix. It is fundamental in ML for applying weights to inputs.

Example: If $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $\mathbf{x} = \begin{bmatrix} 5 \\ 6 \end{bmatrix}$, then $\mathbf{Ax} = \begin{bmatrix} 17 \\ 39 \end{bmatrix}$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
x = np.array([5, 6])
result = np.dot(A, x)
```

1.10. Matrix Multiplication

$$\mathbf{C} = \mathbf{AB}, \quad c_{ij} = \sum_{k=1}^{n} a_{ik} b_{kj}$$

Explanation: Matrix multiplication combines two matrices, producing a matrix that represents the composition of linear transformations. It is used in ML for layer operations in neural networks.

Example: If $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $\mathbf{B} = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$, then $\mathbf{AB} = \begin{bmatrix} 19 & 22 \\ 43 & 50 \end{bmatrix}$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
B = np.array([[5, 6], [7, 8]])
result = np.dot(A, B)
```

Nabin Katwal

1.11.   Transpose of a Matrix

$$\mathbf{A}^T = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}^T = \begin{bmatrix} a_{11} & a_{21} \\ a_{12} & a_{22} \end{bmatrix}$$

Explanation: The transpose of a matrix flips it over its diagonal, exchanging rows with columns. It is used in ML for switching between data representations.

Example: If $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$, then $\mathbf{A}^T = \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix}$.

Implementation:

```python
import numpy as np
A = np.array([[1, 2], [3, 4]])
result = A.T
```

1.12.   Determinant of a 2×2 Matrix

$$\det(\mathbf{A}) = \begin{vmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{vmatrix} = a_{11}a_{22} - a_{12}a_{21}$$

Explanation: The determinant measures the scaling factor of the transformation represented by a matrix. It is used to determine matrix invertibility.

Example: If $\mathbf{A} = \begin{bmatrix} 3 & 8 \\ 4 & 6 \end{bmatrix}$, then $\det(\mathbf{A}) = 3 \cdot 6 - 8 \cdot 4 = -14$.

Implementation:

```python
import numpy as np
A = np.array([[3, 8], [4, 6]])
result = np.linalg.det(A)
```

Nabin Katwal   24

1.13.   Inverse of a 2×2 Matrix

$$\mathbf{A}^{-1} = \frac{1}{\det(\mathbf{A})} \begin{bmatrix} a_{22} & -a_{12} \\ -a_{21} & a_{11} \end{bmatrix}, \quad \det(\mathbf{A}) \neq 0$$

Explanation: The inverse of a 2×2 matrix reverses the linear transformation it represents. It is used in solving systems of linear equations.

Example: If $\mathbf{A} = \begin{bmatrix} 3 & 8 \\ 4 & 6 \end{bmatrix}$, then $\det(\mathbf{A}) = -14$ and $\mathbf{A}^{-1} = \frac{1}{-14} \begin{bmatrix} 6 & -8 \\ -4 & 3 \end{bmatrix}$.

Implementation:

```
import numpy as np
A = np.array([[3, 8], [4, 6]])
result = np.linalg.inv(A)
```

Nabin Katwal

1.14. Cramer's Rule

$$x_i = \frac{\det(\mathbf{A}_i)}{\det(\mathbf{A})}, \quad \det(\mathbf{A}) \neq 0$$

Explanation: Cramer's Rule solves a system of linear equations $\mathbf{Ax} = \mathbf{b}$ by replacing each column of $\mathbf{A}$ with $\mathbf{b}$ and computing determinants. It is a theoretical method often used for small systems.

Example: For $\mathbf{A} = \begin{bmatrix} 2 & 1 \\ 1 & 3 \end{bmatrix}$ and $\mathbf{b} = \begin{bmatrix} 5 \\ 7 \end{bmatrix}$,

$$\mathbf{A}_1 = \begin{bmatrix} 5 & 1 \\ 7 & 3 \end{bmatrix}, \quad \mathbf{A}_2 = \begin{bmatrix} 2 & 5 \\ 1 & 7 \end{bmatrix}$$

and $\det(\mathbf{A}) = 5$, so $x_1 = \frac{\det(\mathbf{A}_1)}{\det(\mathbf{A})}, x_2 = \frac{\det(\mathbf{A}_2)}{\det(\mathbf{A})}$.

Implementation:

```
import numpy as np
A = np.array([[2, 1], [1, 3]])
b = np.array([5, 7])
det_A = np.linalg.det(A)
x = [np.linalg.det(np.column_stack((b if i == j else A[:, j]
                            for j in range(A.shape[1])))) / det_A
     for i in range(A.shape[1])]
```

1.15. Inverse of a Square Matrix

$$\mathbf{A}^{-1} = \frac{1}{\det(\mathbf{A})}\text{adj}(\mathbf{A}), \quad \det(\mathbf{A}) \neq 0$$

Explanation: The inverse of a square matrix generalizes the process for higher dimensions using the adjugate and determinant. It is crucial in linear algebra and ML for solving systems of equations.

Example: If $\mathbf{A} = \begin{bmatrix} 4 & 7 \\ 2 & 6 \end{bmatrix}$, the inverse is computed using cofactor expansion and scaling.

Implementation:

```
import numpy as np
A = np.array([[4, 7], [2, 6]])
result = np.linalg.inv(A)
```

1.16.   Determinant of a Triangular Matrix

$$\det(\mathbf{A}) = \prod_{i=1}^{n} a_{ii}$$

Explanation: The determinant of a triangular matrix (upper or lower) is the product of its diagonal elements. This simplifies determinant calculations and is useful in decompositions.

Example: If $\mathbf{A} = \begin{bmatrix} 2 & 1 & 0 \\ 0 & 3 & 4 \\ 0 & 0 & 5 \end{bmatrix}$, then $\det(\mathbf{A}) = 2 \cdot 3 \cdot 5 = 30$.

Implementation:

```
import numpy as np
A = np.array([[2, 1, 0], [0, 3, 4], [0, 0, 5]])
result = np.prod(np.diag(A))
```

Nabin Katwal

1.17.   Rank-Nullity Theorem

$$\text{rank}(\mathbf{A}) + \text{nullity}(\mathbf{A}) = n$$

Explanation: The Rank-Nullity Theorem states that the sum of the rank (dimension of column space) and nullity (dimension of null space) of a matrix equals the number of columns. It is fundamental in linear algebra for understanding solutions to systems of linear equations.

Example: If $\mathbf{A}$ has 3 columns and its rank is 2, then the nullity is 1 since $2 + 1 = 3$.

Implementation:

```
import numpy as np
from numpy.linalg import matrix_rank
A = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
rank = matrix_rank(A)
nullity = A.shape[1] - rank
```

Nabin Katwal

1.18.   Hadamard (Elementwise) Product

$$C = A \circ B = \begin{bmatrix} a_{11}b_{11} & a_{12}b_{12} \\ a_{21}b_{21} & a_{22}b_{22} \end{bmatrix}$$

Explanation: The Hadamard product performs elementwise multiplication between two matrices. It is used in ML for feature-wise scaling or gating.

Example: If $A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $B = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$, then $C = \begin{bmatrix} 5 & 12 \\ 21 & 32 \end{bmatrix}$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
B = np.array([[5, 6], [7, 8]])
result = np.multiply(A, B)
```

1.19.   Outer Product

$$
\mathbf{C} = \mathbf{u} \otimes \mathbf{v} =
\begin{bmatrix}
u_1v_1 & u_1v_2 & \cdots & u_1v_n \\
u_2v_1 & u_2v_2 & \cdots & u_2v_n \\
\vdots & \vdots & \ddots & \vdots \\
u_mv_1 & u_mv_2 & \cdots & u_mv_n
\end{bmatrix}
$$

Explanation: The outer product generates a matrix by multiplying every element of one vector by every element of another. It is used in tensor operations and constructing rank-1 matrices.

Example: If $\mathbf{u} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$ and $\mathbf{v} = \begin{bmatrix} 3 \\ 4 \\ 5 \end{bmatrix}$, then $\mathbf{u} \otimes \mathbf{v} = \begin{bmatrix} 3 & 4 & 5 \\ 6 & 8 & 10 \end{bmatrix}$.

Implementation:

```
import numpy as np
u = np.array([1, 2])
v = np.array([3, 4, 5])
result = np.outer(u, v)
```

1.20. Frobenius Norm

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i=1}^{m}\sum_{j=1}^{n}|a_{ij}|^2}$$

Explanation: The Frobenius norm measures the magnitude of a matrix by summing the squares of all its elements. It is widely used in optimization and matrix analysis.

Example: If $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$, then $\|\mathbf{A}\|_F = \sqrt{1^2 + 2^2 + 3^2 + 4^2} = \sqrt{30}$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
result = np.linalg.norm(A, 'fro')
```

Nabin Katwal

1.21.   Matrix Norm Inequality

$$\|\mathbf{Ax}\| \leq \|\mathbf{A}\|\|\mathbf{x}\|$$

Explanation: The matrix norm inequality states that the norm of a matrix-vector product is bounded by the product of the matrix norm and the vector norm. It is a key property in numerical linear algebra and ML for error analysis.

Example: For $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $\mathbf{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$, compute $\|\mathbf{Ax}\| \leq \|\mathbf{A}\|\|\mathbf{x}\|$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
x = np.array([1, 1])
left = np.linalg.norm(np.dot(A, x))
right = np.linalg.norm(A) * np.linalg.norm(x)
inequality_holds = left <= right
```

Nabin Katwal   33

1.22. Matrix Trace

$$\text{Tr}(\mathbf{A}) = \sum_{i=1}^{n} a_{ii}$$

Explanation: The trace of a matrix is the sum of its diagonal elements. It is used in ML for loss functions and characterizing matrix properties.

Example: If $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$, then $\text{Tr}(\mathbf{A}) = 1 + 4 = 5$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
result = np.trace(A)
```

1.23. Trace of a Product

$$\mathrm{Tr}(\mathbf{AB}) = \mathrm{Tr}(\mathbf{BA})$$

Explanation: The trace of a product of two matrices is invariant under cyclic permutations. This property is useful in ML for simplifying gradients in matrix calculus.

Example: For $\mathbf{A} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $\mathbf{B} = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$, compute $\mathrm{Tr}(\mathbf{AB}) = \mathrm{Tr}(\mathbf{BA})$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
B = np.array([[5, 6], [7, 8]])
trace1 = np.trace(np.dot(A, B))
trace2 = np.trace(np.dot(B, A))
equality_holds = trace1 == trace2
```

1.24.   Block Matrix Multiplication

$$\mathbf{C} = \begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix} \begin{bmatrix} \mathbf{E} & \mathbf{F} \\ \mathbf{G} & \mathbf{H} \end{bmatrix} = \begin{bmatrix} \mathbf{AE + BG} & \mathbf{AF + BH} \\ \mathbf{CE + DG} & \mathbf{CF + DH} \end{bmatrix}$$

Explanation: Block matrix multiplication follows the same rules as scalar matrix multiplication, but each element is a submatrix. It is used in ML for large-scale computations and decompositions.

Example: Compute the block product for two partitioned 4×4 matrices.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
B = np.array([[5, 6], [7, 8]])
C = np.array([[9, 10], [11, 12]])
D = np.array([[13, 14], [15, 16]])
E = np.array([[17, 18], [19, 20]])
F = np.array([[21, 22], [23, 24]])
G = np.array([[25, 26], [27, 28]])
H = np.array([[29, 30], [31, 32]])
top_left = np.dot(A, E) + np.dot(B, G)
top_right = np.dot(A, F) + np.dot(B, H)
bottom_left = np.dot(C, E) + np.dot(D, G)
bottom_right = np.dot(C, F) + np.dot(D, H)
result = np.block([[top_left, top_right], [bottom_left, bottom_right]])
```

1.25. Kronecker Product

$$C = A \otimes B = \begin{bmatrix} a_{11}B & a_{12}B \\ a_{21}B & a_{22}B \end{bmatrix}$$

Explanation: The Kronecker product produces a block matrix by multiplying each element of one matrix by the entirety of another. It is used in ML for tensor operations and signal processing.

Example: If $A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $B = \begin{bmatrix} 0 & 5 \\ 6 & 7 \end{bmatrix}$, compute $A \otimes B$.

Implementation:

```
import numpy as np
A = np.array([[1, 2], [3, 4]])
B = np.array([[0, 5], [6, 7]])
result = np.kron(A, B)
```

Chapter 2

PROBABILITY AND STATISTICS

2.1.   Conditional Probability

$$P(A \mid B) = \frac{P(A \cap B)}{P(B)}, \quad P(B) > 0$$

Explanation: Conditional probability quantifies the likelihood of event A occurring given that event B has occurred. It is fundamental in probabilistic reasoning and Bayesian inference.

Example: If $P(A \cap B) = 0.2$ and $P(B) = 0.5$, then $P(A \mid B) = \frac{0.2}{0.5} = 0.4$.

Implementation:

```
P_A_and_B = 0.2
P_B = 0.5
P_A_given_B = P_A_and_B / P_B
```

2.2.   Law of Total Probability

$$P(A) = \sum_i P(A \mid B_i)P(B_i)$$

Explanation: The law of total probability relates the probability of an event A to the probabilities of A given a partition of events $\{B_i\}$. It is used in scenarios with conditional dependencies.

Example: If $P(A \mid B_1) = 0.3$, $P(A \mid B_2) = 0.7$, $P(B_1) = 0.4$, and $P(B_2) = 0.6$, then $P(A) = 0.3 \cdot 0.4 + 0.7 \cdot 0.6 = 0.54$.

Implementation:

```
P_A_given_B1 = 0.3
P_A_given_B2 = 0.7
P_B1 = 0.4
P_B2 = 0.6
P_A = P_A_given_B1 * P_B1 + P_A_given_B2 * P_B2
```

2.3. Bayes' Theorem

$$P(A \mid B) = \frac{P(B \mid A)P(A)}{P(B)}, \quad P(B) > 0$$

Explanation: Bayes' Theorem allows the reversal of conditional probabilities, often used in updating beliefs with new evidence in ML and statistics.

Example: If $P(B \mid A) = 0.8$, $P(A) = 0.3$, and $P(B) = 0.5$, then $P(A \mid B) = \frac{0.8 \cdot 0.3}{0.5} = 0.48$.

Implementation:

```
P_B_given_A = 0.8
P_A = 0.3
P_B = 0.5
P_A_given_B = (P_B_given_A * P_A) / P_B
```

2.4. Expectation

$$\mathbb{E}[X] = \sum_i x_i P(X = x_i)$$

Explanation: The expectation (mean) of a random variable is the weighted average of all possible values, weighted by their probabilities. It is central in probability and statistics.

Example: If $X = \{1, 2, 3\}$ with $P(X = 1) = 0.2$, $P(X = 2) = 0.5$, and $P(X = 3) = 0.3$, then $\mathbb{E}[X] = 1 \cdot 0.2 + 2 \cdot 0.5 + 3 \cdot 0.3 = 2.1$.

Implementation:

```
X = [1, 2, 3]
P_X = [0.2, 0.5, 0.3]
expectation = sum(x * p for x, p in zip(X, P_X))
```

Nabin Katwal

2.5. Variance

$$\mathrm{Var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$$

Explanation: Variance measures the spread of a random variable around its mean. It is widely used in ML for assessing uncertainty and model performance.

Example: For $X = \{1, 2, 3\}$ with $P(X = 1) = 0.2$, $P(X = 2) = 0.5$, and $P(X = 3) = 0.3$, compute $\mathbb{E}[X] = 2.1$ and $\mathbb{E}[X^2] = 4.7$, so $\mathrm{Var}(X) = 4.7 - (2.1)^2 = 0.29$.

Implementation:

```
X = [1, 2, 3]
P_X = [0.2, 0.5, 0.3]
expectation = sum(x * p for x, p in zip(X, P_X))
expectation_X2 = sum(x**2 * p for x, p in zip(X, P_X))
variance = expectation_X2 - expectation**2
```

Nabin Katwal

2.6. Standard Deviation

$$\sigma(X) = \sqrt{\mathrm{Var}(X)}$$

Explanation: The standard deviation is the square root of the variance and provides a measure of dispersion in the same units as the random variable. It is widely used in data analysis and ML for variability assessment.

Example: If $\mathrm{Var}(X) = 0.29$, then $\sigma(X) = \sqrt{0.29} \approx 0.54$.

Implementation:

```
variance = 0.29
std_dev = variance**0.5
```

2.7.   Covariance

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$$

Explanation: Covariance measures the joint variability of two random variables. A positive value indicates that they increase together, while a negative value indicates an inverse relationship.

Example: If $X = \{1, 2\}$, $Y = \{3, 4\}$, $P(X, Y) = \{0.5, 0.5\}$, and $\mathbb{E}[X] = 1.5$, $\mathbb{E}[Y] = 3.5$, compute $\text{Cov}(X, Y) = 0.25$.

Implementation:

```
X = [1, 2]
Y = [3, 4]
P_XY = [0.5, 0.5]
E_X = sum(x * p for x, p in zip(X, P_XY))
E_Y = sum(y * p for y, p in zip(Y, P_XY))
covariance = sum((x - E_X) * (y - E_Y) * p for x, y, p in zip(X, Y, P_XY))
```

2.8.   Correlation

$$\rho(X, Y) = \frac{\text{Cov}(X, Y)}{\sigma(X)\sigma(Y)}$$

Explanation: Correlation normalizes covariance to a scale of $[-1, 1]$, quantifying the strength and direction of a linear relationship between two variables.

Example: If $\text{Cov}(X, Y) = 0.25$, $\sigma(X) = 0.5$, and $\sigma(Y) = 1.0$, then $\rho(X, Y) = \frac{0.25}{0.5 \cdot 1.0} = 0.5$.

Implementation:

```
covariance = 0.25
std_X = 0.5
std_Y = 1.0
correlation = covariance / (std_X * std_Y)
```

2.9.   Probability Mass Function (PMF)

$$P(X = x) = \begin{cases} p_i, & \text{if } x = x_i \\ 0, & \text{otherwise} \end{cases}$$

Explanation: The PMF defines the probabilities of discrete outcomes of a random variable. It is a foundational concept in probability theory.

Example: If $X = \{1, 2, 3\}$ with $P(X = 1) = 0.2$, $P(X = 2) = 0.5$, and $P(X = 3) = 0.3$, the PMF is defined for these values.

Implementation:

```python
X = [1, 2, 3]
P_X = [0.2, 0.5, 0.3]
def pmf(x):
    return P_X[X.index(x)] if x in X else 0
```

Nabin Katwal

2.10.   Probability Density Function (PDF)

$$f_X(x) \geq 0, \quad \int_{-\infty}^{\infty} f_X(x)dx = 1$$

Explanation: The PDF defines the relative likelihood of a continuous random variable at a specific value. It is used in probability and statistics for modeling continuous distributions.

Example: For a standard normal distribution, the PDF is $f_X(x) = \frac{1}{\sqrt{2\pi}}e^{-\frac{x^2}{2}}$.

Implementation:

```
import numpy as np
from scipy.stats import norm
x = 0  # example point
pdf_value = norm.pdf(x)
```

2.11.   Joint Probability

$$P(A \cap B) = P(A \mid B)P(B)$$

Explanation: Joint probability quantifies the likelihood of two events occurring together. It is essential in probabilistic modeling and understanding relationships between variables.

Example: If $P(A \mid B) = 0.4$ and $P(B) = 0.5$, then $P(A \cap B) = 0.4 \cdot 0.5 = 0.2$.

Implementation:

```
P_A_given_B = 0.4
P_B = 0.5
P_A_and_B = P_A_given_B * P_B
```

Nabin Katwal   49

2.12.   CDF (Cumulative Distribution Function)

$$F_X(x) = P(X \leq x)$$

Explanation: The CDF of a random variable gives the probability that the variable takes a value less than or equal to x. It is used to describe the distribution function for both discrete and continuous variables.

Example: For a uniform distribution $X \sim U(0, 1)$, $F_X(0.5) = 0.5$.

Implementation:

```
from scipy.stats import uniform
x = 0.5
cdf_value = uniform.cdf(x, loc=0, scale=1)
```

2.13. Entropy (discrete)

$$H(X) = -\sum_i P(X = x_i) \log_2 P(X = x_i)$$

Explanation: Entropy measures the uncertainty of a discrete random variable. It is a fundamental concept in information theory and ML, particularly in decision trees and loss functions.

Example: If $P(X) = \{0.5, 0.5\}$, then $H(X) = -0.5 \log_2(0.5) - 0.5 \log_2(0.5) = 1$.

Implementation:

```
import numpy as np
P_X = [0.5, 0.5]
entropy = -sum(p * np.log2(p) for p in P_X if p > 0)
```

2.14.   Conditional Expectation

$$\mathbb{E}[X \mid Y] = \sum_x x P(X = x \mid Y)$$

Explanation: Conditional expectation is the expected value of a random variable X given that another variable Y is known. It is critical in Bayesian inference and probabilistic modeling.

Example: If $X = \{1, 2\}$ with $P(X = 1 \mid Y) = 0.7$ and $P(X = 2 \mid Y) = 0.3$, then $\mathbb{E}[X \mid Y] = 1 \cdot 0.7 + 2 \cdot 0.3 = 1.3$.

Implementation:

```
X = [1, 2]
P_X_given_Y = [0.7, 0.3]
conditional_expectation = sum(x * p for x, p in zip(X, P_X_given_Y))
```

2.15. Law of Iterated Expectations

$$\mathbb{E}[X] = \mathbb{E}[\mathbb{E}[X \mid Y]]$$

Explanation: The law of iterated expectations states that the expectation of X is the weighted average of its conditional expectations over Y. It is foundational in probability theory and statistics.

Example: Suppose X depends on $Y = \{1, 2\}$, with $\mathbb{E}[X \mid Y = 1] = 3$, $\mathbb{E}[X \mid Y = 2] = 5$, and $P(Y = 1) = 0.6$, $P(Y = 2) = 0.4$. Then $\mathbb{E}[X] = 3 \cdot 0.6 + 5 \cdot 0.4 = 3.8$.

Implementation:

```
E_X_given_Y = [3, 5]
P_Y = [0.6, 0.4]
E_X = sum(e * p for e, p in zip(E_X_given_Y, P_Y))
```

Nabin Katwal

2.16. Marginal Probability

$$P(A) = \sum_B P(A \cap B)$$

Explanation: Marginal probability calculates the probability of an event A by summing (or integrating, for continuous cases) over all possible outcomes of another variable B. It is used in probabilistic modeling to reduce joint distributions.

Example: If $P(A \cap B_1) = 0.3$ and $P(A \cap B_2) = 0.4$, then $P(A) = 0.3 + 0.4 = 0.7$.

Implementation:

```
P_A_and_B = [0.3, 0.4]
P_A = sum(P_A_and_B)
```

2.17. Skewness

$$\text{Skewness}(X) = \frac{\mathbb{E}[(X - \mu)^3]}{\sigma^3}$$

Explanation: Skewness measures the asymmetry of the probability distribution of a random variable about its mean. Positive skew indicates a longer right tail, and negative skew indicates a longer left tail.

Example: For $X = \{1, 2, 3\}$ with mean $\mu = 2$ and standard deviation $\sigma = 0.816$, compute Skewness(X) using the third central moment.

Implementation:

```
import numpy as np
X = [1, 2, 3]
mu = np.mean(X)
sigma = np.std(X)
skewness = np.mean(((X - mu) / sigma)**3)
```

2.18.   Kurtosis

$$\text{Kurtosis}(X) = \frac{\mathbb{E}[(X - \mu)^4]}{\sigma^4}$$

Explanation: Kurtosis measures the "tailedness" of the probability distribution. A high kurtosis indicates heavy tails, while a low kurtosis indicates light tails.

Example: For $X = \{1, 2, 3\}$ with mean $\mu = 2$ and standard deviation $\sigma = 0.816$, compute Kurtosis(X) using the fourth central moment.

Implementation:

```
import numpy as np
X = [1, 2, 3]
mu = np.mean(X)
sigma = np.std(X)
kurtosis = np.mean(((X - mu) / sigma)**4)
```

Nabin Katwal

2.19.   Binary Cross-Entropy (special case)

$$\text{BCE}(y, \hat{y}) = -\frac{1}{n} \sum_{i=1}^{n} (y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i))$$

Explanation: Binary cross-entropy is a loss function used for binary classification tasks. It measures the dissimilarity between predicted probabilities $(\hat{y})$ and true labels (y).

Example: For $y = [1, 0]$ and $\hat{y} = [0.8, 0.2]$, compute $\text{BCE} = -\frac{1}{2} (\log(0.8) + \log(0.8))$.

Implementation:

```
import numpy as np
y = np.array([1, 0])
y_hat = np.array([0.8, 0.2])
bce = -np.mean(y * np.log(y_hat) + (1 - y) * np.log(1 - y_hat))
```

Nabin Katwal

2.20. Variance (Alternative)

$$\mathrm{Var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$$

Explanation: An alternative formula for variance uses the difference between the expected value of the square of X and the square of the expected value of X. This method is computationally efficient.

Example: For $X = \{1, 2, 3\}$, compute $\mathbb{E}[X^2] = \frac{1^2 + 2^2 + 3^2}{3} = 4.67$ and $(\mathbb{E}[X])^2 = 2^2 = 4$, so $\mathrm{Var}(X) = 0.67$.

Implementation:

```
import numpy as np
X = np.array([1, 2, 3])
E_X2 = np.mean(X**2)
E_X = np.mean(X)
variance = E_X2 - E_X**2
```

Chapter 3

CALCULUS

3.1.   Limit Definition of Derivative

$$f'(x) = \lim_{h \to 0} \frac{f(x + h) - f(x)}{h}$$

Explanation: The derivative of a function is defined as the limit of the difference quotient as h approaches zero. It represents the instantaneous rate of change of the function.

Example: For $f(x) = x^2$, compute $f'(x) = \lim_{h \to 0} \frac{(x+h)^2 - x^2}{h} = 2x$.

Implementation:

```
def derivative(f, x, h=1e-5):
    return (f(x + h) - f(x)) / h
```

3.2. Power Rule

$$\frac{d}{dx}x^n = nx^{n-1}$$

Explanation: The power rule simplifies differentiation of monomials. It is foundational for calculus and widely used in gradient computations in ML.

Example: For $f(x) = x^3$, $f'(x) = 3x^2$.

Implementation:

```
def power_rule(n, x):
    return n * x**(n - 1)
```

3.3. Product Rule

$$\frac{d}{dx}[u(x)v(x)] = u'(x)v(x) + u(x)v'(x)$$

Explanation: The product rule computes the derivative of the product of two functions. It is crucial for handling multiplicative relationships in ML.

Example: For $f(x) = (x^2)(e^x)$, $f'(x) = 2xe^x + x^2e^x$.

Implementation:

```
def product_rule(u, v, u_prime, v_prime, x):
    return u_prime(x) * v(x) + u(x) * v_prime(x)
```

Nabin Katwal

3.4.   Quotient Rule

$$\frac{d}{dx}\left[\frac{u(x)}{v(x)}\right] = \frac{u'(x)v(x) - u(x)v'(x)}{[v(x)]^2}$$

Explanation: The quotient rule computes the derivative of the ratio of two functions. It is essential for operations involving divisions in ML models.

Example: For $f(x) = \frac{x^2}{e^x}$, $f'(x) = \frac{2xe^x - x^2 e^x}{e^{2x}}$.

Implementation:

```
def quotient_rule(u, v, u_prime, v_prime, x):
    return (u_prime(x) * v(x) - u(x) * v_prime(x)) / (v(x)**2)
```

3.5. Chain Rule

$$\frac{d}{dx}f(g(x)) = f'(g(x))g'(x)$$

Explanation: The chain rule computes the derivative of a composite function. It is extensively used in backpropagation for training neural networks.

Example: For $f(x) = \sin(x^2)$, $f'(x) = \cos(x^2) \cdot 2x$.

Implementation:

```
def chain_rule(f_prime, g, g_prime, x):
    return f_prime(g(x)) * g_prime(x)
```

Nabin Katwal

3.6. Logarithmic Derivative

$$\frac{d}{dx}\ln(x) = \frac{1}{x}, \quad x > 0$$

Explanation: The derivative of the natural logarithm function is the reciprocal of its argument. It is frequently used in ML for optimization and logarithmic transformations.

Example: For $f(x) = \ln(x)$, $f'(2) = \frac{1}{2}$.

Implementation:

```python
import numpy as np
def log_derivative(x):
    return 1 / x
```

3.7. Exponential Derivative

$$\frac{d}{dx}e^x = e^x$$

Explanation: The exponential function is unique as its derivative is equal to itself. This property is key in gradient computations and exponential growth models in ML.

Example: For $f(x) = e^x$, $f'(2) = e^2$.

Implementation:

```python
import numpy as np
def exp_derivative(x):
    return np.exp(x)
```

3.8.   Integral of a Power Function

$$\int x^n dx = \frac{x^{n+1}}{n+1} + C, \quad n \neq -1$$

Explanation: The integral of a power function generalizes the antiderivative for monomials. This rule is fundamental in integral calculus and applied in ML for cost function analysis.

Example: For $f(x) = x^2$, $\int x^2 dx = \frac{x^3}{3} + C$.

Implementation:

```python
def power_integral(n, x):
    return x**(n + 1) / (n + 1)
```

3.9.   Fundamental Theorem of Calculus

$$\int_a^b f(x)dx = F(b) - F(a), \quad \text{where } F'(x) = f(x)$$

Explanation: The Fundamental Theorem of Calculus links differentiation and integration, stating that integration over an interval is the difference of the antiderivative evaluated at the endpoints.

Example: For $f(x) = x^2$ over $[1,3]$, $\int_1^3 x^2 dx = \left[\frac{x^3}{3}\right]_1^3 = \frac{27}{3} - \frac{1}{3} = \frac{26}{3}$.

Implementation:

```
def definite_integral(f, a, b):
    from scipy.integrate import quad
    result, _ = quad(f, a, b)
    return result
```

Nabin Katwal

3.10. Partial Derivatives

$$\frac{\partial f}{\partial x} = \lim_{h \to 0} \frac{f(x + h, y) - f(x, y)}{h}, \quad \frac{\partial f}{\partial y} = \lim_{h \to 0} \frac{f(x, y + h) - f(x, y)}{h}$$

Explanation: Partial derivatives measure the rate of change of a multivariable function with respect to one variable while keeping others constant. They are essential in optimization and gradient-based ML methods.

Example: For $f(x, y) = x^2 + y^2$, $\frac{\partial f}{\partial x} = 2x$, $\frac{\partial f}{\partial y} = 2y$.

Implementation:

```
def partial_derivative(f, var, point, h=1e-5):
    args = list(point)
    args[var] += h
    return (f(*args) - f(*point)) / h
```

Nabin Katwal   68

3.11.   Gradient

$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

Explanation: The gradient is a vector containing all partial derivatives of a scalar-valued function. It points in the direction of the steepest ascent and is widely used in ML optimization algorithms like gradient descent.

Example: For $f(x, y) = x^2 + y^2$, $\nabla f(x, y) = \begin{bmatrix} 2x \\ 2y \end{bmatrix}$.

Implementation:

```
import numpy as np
def gradient(f, point, h=1e-5):
    grad = np.zeros(len(point))
    for i in range(len(point)):
        args = point.copy()
        args[i] += h
        grad[i] = (f(*args) - f(*point)) / h
    return grad
```

3.12.   Second Derivative (Hessian)

$$H(f) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots \\ \vdots & \vdots & \ddots \end{bmatrix}$$

Explanation: The Hessian is a square matrix of second-order partial derivatives. It is used in optimization to assess curvature and convergence properties of a function.

Example: For $f(x, y) = x^2 + y^2$, the Hessian is $H(f) = \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix}$.

Implementation:

```
def hessian(f, point, h=1e-5):
    n = len(point)
    hess = np.zeros((n, n))
    for i in range(n):
        for j in range(n):
            args = point.copy()
            args[i] += h
            args[j] += h
            f_ij = f(*args)
            args[j] -= h
            f_i = f(*args)
            args[i] -= h
            args[j] += h
            f_j = f(*args)
            f_orig = f(*point)
```

```
        hess[i, j] = (f_ij - f_i - f_j + f_orig) / (h ** 2)

    return hess
```

3.13. Directional Derivative

$$D_{\mathbf{v}}f(\mathbf{x}) = \nabla f(\mathbf{x}) \cdot \mathbf{v}$$

Explanation: The directional derivative measures the rate of change of a function in the direction of a given vector. It is critical in optimization and ML for evaluating function behavior in a specific direction.

Example: For $f(x, y) = x^2 + y^2$, $\nabla f(x, y) = \begin{bmatrix} 2x \\ 2y \end{bmatrix}$. In the direction $\mathbf{v} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$, $D_{\mathbf{v}}f(x, y) = 2x$.

Implementation:

```
def directional_derivative(f, grad_f, point, direction):
    grad = grad_f(point)
    return np.dot(grad, direction)
```

3.14.   Higher-Order Partial Derivatives

$$\frac{\partial^k f}{\partial x_1^{p_1} \partial x_2^{p_2} \cdots \partial x_n^{p_n}}$$

Explanation: Higher-order partial derivatives extend partial derivatives to greater orders. Mixed derivatives often satisfy equality $(f_{xy} = f_{yx})$ under smoothness conditions.

Example: For $f(x, y) = x^2 y$, $\frac{\partial^2 f}{\partial x \partial y} = 2x$.

Implementation:

```
def higher_order_partial(f, point, var_indices, h=1e-5):
    args = list(point)
    for var in var_indices:
        args[var] += h
    f_plus = f(*args)
    for var in var_indices:
        args[var] -= h * len(var_indices)
    f_minus = f(*args)
    return (f_plus - f_minus) / (h ** len(var_indices))
```

Nabin Katwal

3.15. Total Derivative

$$\frac{df}{dt} = \sum_{i=1}^{n} \frac{\partial f}{\partial x_i} \frac{dx_i}{dt}$$

Explanation: The total derivative accounts for changes in all independent variables as functions of an external variable t. It is used in dynamical systems and optimization.

Example: If $f(x, y) = x^2 + y^2$, $x = t$, and $y = t^2$, then $\frac{df}{dt} = 2x \cdot 1 + 2y \cdot 2t = 2t + 4t^3$.

Implementation:

```
def total_derivative(f, partials, dx_dt, point):
    return sum(partials[i] * dx_dt[i] for i in range(len(point)))
```

3.16.   Implicit Differentiation

$$\frac{dy}{dx} = -\frac{\frac{\partial F}{\partial x}}{\frac{\partial F}{\partial y}}$$

Explanation: Implicit differentiation computes the derivative of a dependent variable in an equation where the variable cannot be explicitly solved. It is used in ML and calculus for handling complex equations.

Example: For $F(x, y) = x^2 + y^2 - 1 = 0$, $\frac{dy}{dx} = -\frac{x}{y}$.

Implementation:

```
def implicit_differentiation(F, x, y, partial_F_x, partial_F_y):
    return -partial_F_x(x, y) / partial_F_y(x, y)
```

3.17.   Taylor Series Expansion

$$f(x) \approx f(a) + f'(a)(x - a) + \frac{f''(a)}{2!}(x - a)^2 + \cdots$$

Explanation: The Taylor series approximates a function near a point a using its derivatives. It is used in optimization and numerical analysis.

Example: For $f(x) = e^x$ near $a = 0$, $f(x) \approx 1 + x + \frac{x^2}{2} + \cdots$.

Implementation:

```
def taylor_series(f, derivatives, a, x, terms=3):
    result = 0
    for n in range(terms):
        result += derivatives[n](a) * (x - a)**n / np.math.factorial(n)
    return result
```

3.18.   Jacobian Matrix

$$J(\mathbf{f}) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

Explanation: The Jacobian matrix contains all first-order partial derivatives of a vector-valued function. It is essential in ML for gradient-based optimization in multivariable spaces.

Example: For $\mathbf{f}(x, y) = \begin{bmatrix} x^2 + y \\ y^2 + x \end{bmatrix}$, the Jacobian is $\begin{bmatrix} 2x & 1 \\ 1 & 2y \end{bmatrix}$.

Implementation:

```
def jacobian(f, point, h=1e-5):
    m = len(f)
    n = len(point)
    J = np.zeros((m, n))
    for i in range(m):
        for j in range(n):
            args = point.copy()
            args[j] += h
            J[i, j] = (f[i](*args) - f[i](*point)) / h
    return J
```

3.19. Arc Length of a Curve

$$L = \int_a^b \sqrt{1 + \left(\frac{dy}{dx}\right)^2}\, dx$$

Explanation: The arc length measures the distance along a curve between two points. It is used in geometry and physics for path analysis.

Example: For $y = x^2$ over $[0, 1]$, $L = \int_0^1 \sqrt{1 + (2x)^2}\, dx$.

Implementation:

```python
from scipy.integrate import quad
def arc_length(f_prime, a, b):
    integrand = lambda x: np.sqrt(1 + f_prime(x)**2)
    return quad(integrand, a, b)[0]
```

3.20.   Curvature of a Function

$$\kappa(x) = \frac{|y''(x)|}{(1 + [y'(x)]^2)^{3/2}}$$

Explanation: Curvature quantifies how sharply a curve bends at a given point. It is used in geometry and trajectory analysis in robotics and ML.

Example: For $y = x^2$, $y'(x) = 2x$, $y''(x) = 2$, so $\kappa(x) = \frac{2}{(1+4x^2)^{3/2}}$.

Implementation:

```
def curvature(f_prime, f_double_prime, x):
    numerator = abs(f_double_prime(x))
    denominator = (1 + f_prime(x)**2)**1.5
    return numerator / denominator
```

Nabin Katwal

3.21.   Integral by Parts

$$\int uv'dx = uv - \int u'vdx$$

Explanation: Integration by parts is a technique derived from the product rule of differentiation. It is used to simplify integrals involving products of functions.

Example: For $\int xe^x dx$, let $u = x$ and $v' = e^x$. Then $\int xe^x dx = xe^x - \int e^x dx = xe^x - e^x + C$.

Implementation:

```python
from sympy import symbols, integrate, exp
x = symbols('x')
u = x
v_prime = exp(x)
v = integrate(v_prime, x)
integral = u * v - integrate(v * u.diff(x), x)
```

3.22.   Volume of Revolution (Disk Method)

$$V = \pi \int_a^b [f(x)]^2 dx$$

Explanation: The disk method computes the volume of a solid of revolution by slicing it into disks perpendicular to the axis of rotation. It is common in geometry and physics.

Example: For $y = x^2$ revolved around the x-axis over $[0, 1]$, $V = \pi \int_0^1 (x^2)^2 dx = \pi \int_0^1 x^4 dx = \frac{\pi}{5}$.

Implementation:

```python
from scipy.integrate import quad
import numpy as np
def volume_of_revolution(f, a, b):
    integrand = lambda x: np.pi * f(x)**2
    return quad(integrand, a, b)[0]
```

3.23.   Surface Integral

$$\iint_S f(x, y, z)dS = \iint_R f(x, y, g(x, y))\sqrt{1 + \left(\frac{\partial g}{\partial x}\right)^2 + \left(\frac{\partial g}{\partial y}\right)^2}\, dA$$

Explanation: A surface integral extends the idea of a line integral to a surface, summing a scalar field or vector flux over the surface.

Example: Compute the surface integral of $f(x, y, z) = z$ over $z = x^2 + y^2$ for $x^2 + y^2 \leq 1$.

Implementation:

```
from scipy.integrate import dblquad
def surface_integral(f, g, bounds_x, bounds_y):
    def integrand(x, y):
        gx, gy = g(x, y)
        return f(x, y, g(x, y)) * np.sqrt(1 + gx**2 + gy**2)
    return dblquad(integrand, *bounds_x, *bounds_y)
```

3.24.   Divergence of a Vector Field

$$\text{div}\,\mathbf{F} = \nabla \cdot \mathbf{F} = \frac{\partial F_1}{\partial x} + \frac{\partial F_2}{\partial y} + \frac{\partial F_3}{\partial z}$$

Explanation: The divergence measures the magnitude of a vector field's source or sink at a given point. It is used in fluid dynamics and electromagnetism.

Example: For $\mathbf{F} = \begin{bmatrix} x \\ y \\ z \end{bmatrix}$, $\text{div}\,\mathbf{F} = 1 + 1 + 1 = 3$.

Implementation:

```
from sympy import symbols, diff
x, y, z = symbols('x y z')
F = [x, y, z]
divergence = sum(diff(F[i], var) for i, var in enumerate([x, y, z]))
```

Nabin Katwal

3.25.   Curl of a Vector Field

$$\text{curl}\,\mathbf{F} = \nabla \times \mathbf{F} = \begin{vmatrix} \mathbf{i} & \mathbf{j} & \mathbf{k} \\ \frac{\partial}{\partial x} & \frac{\partial}{\partial y} & \frac{\partial}{\partial z} \\ F_1 & F_2 & F_3 \end{vmatrix}$$

Explanation: The curl measures the rotation or circulation of a vector field at a point. It is critical in fluid mechanics and electromagnetism.

Example: For $\mathbf{F} = \begin{bmatrix} 0 \\ 0 \\ xy \end{bmatrix}$, $\text{curl}\,\mathbf{F} = \begin{bmatrix} -y \\ x \\ 0 \end{bmatrix}$.

Implementation:

```
from sympy import symbols, Matrix
x, y, z = symbols('x y z')
F = Matrix([0, 0, x*y])
curl = F.jacobian([x, y, z]).transpose() - F.jacobian([x, y, z])
```

Chapter 4

OPTIMIZATION

4.1. Gradient Descent

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla J(\theta^{(t)})$$

Explanation: Gradient descent is an optimization algorithm that iteratively updates parameters in the direction of the negative gradient to minimize the cost function $J(\theta)$.

Example: For $J(\theta) = \theta^2$ and $\eta = 0.1$, the update is $\theta^{(t+1)} = \theta^{(t)} - 0.2\theta^{(t)}$.

Implementation:

```
def gradient_descent(gradient, theta, eta, steps):
    for _ in range(steps):
        theta -= eta * gradient(theta)
    return theta
```

4.2. Stochastic Gradient Descent (SGD)

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla J_i(\theta^{(t)})$$

Explanation: SGD computes gradients on individual data points, updating parameters more frequently. It is widely used in ML due to its efficiency with large datasets.

Example: For $J_i(\theta) = (\theta - y_i)^2$, the update is based on one data point at each iteration.

Implementation:

```
def stochastic_gradient_descent(gradient, theta, eta, data, steps):
    for _ in range(steps):
        i = np.random.randint(len(data))
        theta -= eta * gradient(theta, data[i])
    return theta
```

4.3.   Momentum-based Gradient Descent

$$v^{(t+1)} = \beta v^{(t)} - \eta \nabla J(\theta^{(t)}), \quad \theta^{(t+1)} = \theta^{(t)} + v^{(t+1)}$$

Explanation: Momentum adds an exponentially weighted moving average of past gradients to the current update, improving convergence speed and stability.

Example: For $\beta = 0.9$, $\eta = 0.1$, the velocity update smooths oscillations in gradient descent.

Implementation:

```
def momentum_gradient_descent(gradient, theta, eta, beta, steps):
    v = 0
    for _ in range(steps):
        v = beta * v - eta * gradient(theta)
        theta += v
    return theta
```

Nabin Katwal

4.4. Nesterov Accelerated Gradient (NAG)

$$v^{(t+1)} = \beta v^{(t)} - \eta \nabla J(\theta^{(t)} + \beta v^{(t)}), \quad \theta^{(t+1)} = \theta^{(t)} + v^{(t+1)}$$

Explanation: NAG improves upon momentum by calculating gradients at a lookahead position, resulting in more precise updates.

Example: For $\beta = 0.9$, NAG anticipates the future direction, reducing overshooting in oscillatory scenarios.

Implementation:

```
def nesterov_gradient_descent(gradient, theta, eta, beta, steps):
    v = 0
    for _ in range(steps):
        lookahead = theta + beta * v
        v = beta * v - eta * gradient(lookahead)
        theta += v
    return theta
```

4.5. RMSProp

$$s^{(t+1)} = \beta s^{(t)} + (1 - \beta)[\nabla J(\theta^{(t)})]^2, \quad \theta^{(t+1)} = \theta^{(t)} - \frac{\eta}{\sqrt{s^{(t+1)} + \epsilon}} \nabla J(\theta^{(t)})$$

Explanation: RMSProp scales the learning rate by a moving average of squared gradients, improving convergence for non-convex problems.

Example: For $\beta = 0.9$, RMSProp adapts the step size for each parameter, stabilizing updates.

Implementation:

```
def rmsprop(gradient, theta, eta, beta, epsilon, steps):
    s = 0
    for _ in range(steps):
        grad = gradient(theta)
        s = beta * s + (1 - beta) * grad**2
        theta -= eta / (np.sqrt(s) + epsilon) * grad
    return theta
```

4.6.   Adam Optimization

$$m^{(t+1)} = \beta_1 m^{(t)} + (1 - \beta_1)\nabla J(\theta^{(t)}), \quad s^{(t+1)} = \beta_2 s^{(t)} + (1 - \beta_2)[\nabla J(\theta^{(t)})]^2$$

$$\hat{m} = \frac{m^{(t+1)}}{1 - \beta_1^t}, \quad \hat{s} = \frac{s^{(t+1)}}{1 - \beta_2^t}, \quad \theta^{(t+1)} = \theta^{(t)} - \frac{\eta}{\sqrt{\hat{s}} + \epsilon}\hat{m}$$

Explanation: Adam combines momentum and RMSProp, adapting step sizes and smoothing updates. It is one of the most popular optimization algorithms in ML.

Example: For $\beta_1 = 0.9$, $\beta_2 = 0.999$, Adam balances momentum and per-parameter scaling.

Implementation:

```
def adam(gradient, theta, eta, beta1, beta2, epsilon, steps):
    m, s = 0, 0
    for t in range(1, steps + 1):
        grad = gradient(theta)
        m = beta1 * m + (1 - beta1) * grad
        s = beta2 * s + (1 - beta2) * grad**2
        m_hat = m / (1 - beta1**t)
        s_hat = s / (1 - beta2**t)
        theta -= eta / (np.sqrt(s_hat) + epsilon) * m_hat
    return theta
```

4.7.   Regularized Optimization Objective

$$J_{\text{reg}}(\theta) = J(\theta) + \lambda R(\theta)$$

Explanation: Regularization penalizes model complexity to prevent overfitting. Common regularizers include L1 (lasso) and L2 (ridge) norms.

Example: For $R(\theta) = \|\theta\|_2^2$, $J_{\text{reg}}(\theta) = J(\theta) + \lambda\|\theta\|_2^2$.

Implementation:

```
def regularized_objective(loss, theta, reg, lam):
    return loss(theta) + lam * reg(theta)
```

4.8. Learning Rate Decay

$$\eta_t = \frac{\eta_0}{1 + \gamma t}$$

Explanation: Learning rate decay gradually reduces the learning rate to improve convergence stability as training progresses.

Example: For $\eta_0 = 0.1$, $\gamma = 0.01$, at step $t = 10$, $\eta_t = 0.1/(1 + 0.01 \cdot 10) = 0.0909$.

Implementation:

```
def learning_rate_decay(eta0, gamma, t):
    return eta0 / (1 + gamma * t)
```

4.9. Gradient Clipping

$$\mathbf{g} = \text{clip}(\mathbf{g}, -\tau, \tau)$$

Explanation: Gradient clipping limits the gradient magnitude to prevent exploding gradients in deep neural networks.

Example: For $\tau = 1.0$, clip gradients to the range $[-1, 1]$.

Implementation:

```
def gradient_clipping(grad, tau):
    return np.clip(grad, -tau, tau)
```

4.10.   Minibatch Gradient Descent

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla J_{B_t}(\theta^{(t)})$$

Explanation: Minibatch gradient descent computes updates using small random subsets of data, balancing SGD's noise and batch gradient descent's stability.

Example: Use minibatch size $B = 32$ to compute updates on smaller subsets of data.

Implementation:

```
def minibatch_gradient_descent(gradient, theta, eta, data, batch_size, steps):
    for _ in range(steps):
        batch = np.random.choice(data, batch_size, replace=False)
        theta -= eta * gradient(theta, batch)
    return theta
```

4.11.   Coordinate Descent

$$\theta_j^{(t+1)} = \theta_j^{(t)} - \eta \frac{\partial J(\theta)}{\partial \theta_j}$$

Explanation: Coordinate descent optimizes a single parameter at a time, cycling through all parameters until convergence. It is effective for high-dimensional problems.

Example: Minimize $J(\theta_1, \theta_2) = (\theta_1 - 1)^2 + (\theta_2 - 2)^2$ by alternately updating θ_1 and θ_2.

Implementation:

```
def coordinate_descent(gradient, theta, eta, steps):
    for _ in range(steps):
        for j in range(len(theta)):
            theta[j] -= eta * gradient(theta, j)
    return theta
```

Nabin Katwal   95

4.12.   Elastic Net Regularization

$$J_{\text{reg}}(\theta) = J(\theta) + \lambda_1 \|\theta\|_1 + \lambda_2 \|\theta\|_2^2$$

Explanation: Elastic Net combines L1 and L2 regularization to handle sparsity and multicollinearity. It is commonly used in regression tasks.

Example: For $\lambda_1 = 0.1$, $\lambda_2 = 0.2$, and $J(\theta) = \|\theta - \mathbf{y}\|_2^2$, compute the regularized objective.

Implementation:

```
def elastic_net_objective(loss, theta, lam1, lam2):
    return loss(theta) + lam1 * np.sum(np.abs(theta)) + lam2 * np.sum(theta**2)
```

4.13.   Adagrad Optimization

$$\theta^{(t+1)} = \theta^{(t)} - \frac{\eta}{\sqrt{G^{(t)} + \epsilon}} \nabla J(\theta^{(t)})$$

$$G^{(t)} = \sum_{i=1}^{t} [\nabla J(\theta^{(i)})]^2$$

Explanation: Adagrad adapts the learning rate for each parameter based on the history of gradients, improving performance on sparse data.

Example: For $\eta = 0.1$, adaptively scale updates for different features.

Implementation:

```
def adagrad(gradient, theta, eta, epsilon, steps):
    G = 0
    for _ in range(steps):
        grad = gradient(theta)
        G += grad**2
        theta -= eta / (np.sqrt(G) + epsilon) * grad
    return theta
```

Nabin Katwal

4.14.   AdamW Optimization

$$\theta^{(t+1)} = \theta^{(t)} - \frac{\eta}{\sqrt{\hat{s}} + \epsilon}\hat{m} - \lambda\theta^{(t)}$$

Explanation: AdamW modifies Adam by decoupling weight decay from the gradient updates, improving regularization and generalization in ML models.

Example: For $\lambda = 0.01$, regularize weights alongside adaptive learning rates.

Implementation:

```
def adamw(gradient, theta, eta, beta1, beta2, lam, epsilon, steps):
    m, s = 0, 0
    for t in range(1, steps + 1):
        grad = gradient(theta)
        m = beta1 * m + (1 - beta1) * grad
        s = beta2 * s + (1 - beta2) * grad**2
        m_hat = m / (1 - beta1**t)
        s_hat = s / (1 - beta2**t)
        theta -= eta / (np.sqrt(s_hat) + epsilon) * m_hat + lam * theta
    return theta
```

4.15. Momentum "Heavy Ball" Method

$$\theta^{(t+1)} = \theta^{(t)} + \beta(\theta^{(t)} - \theta^{(t-1)}) - \eta\nabla J(\theta^{(t)})$$

Explanation: This variant of momentum includes an inertial term to improve convergence speed for strongly convex problems.

Example: For $\beta = 0.9$, the "heavy ball" accelerates gradient descent.

Implementation:

```
def heavy_ball(gradient, theta, eta, beta, steps):
    prev_theta = theta.copy()
    v = 0
    for _ in range(steps):
        grad = gradient(theta)
        v = beta * (theta - prev_theta) - eta * grad
        prev_theta = theta.copy()
        theta += v
    return theta
```

4.16. Projection / Projected Gradient Descent

$$\theta^{(t+1)} = \text{Proj}_{\mathcal{C}}(\theta^{(t)} - \eta \nabla J(\theta^{(t)}))$$

Explanation: Projected gradient descent ensures that updates remain within a feasible set $\mathcal{C}$, often used for constrained optimization.

Example: For $\mathcal{C} = \|\theta\|_2 \leq 1$, project θ onto the unit ball after each step.

Implementation:

```python
def projected_gradient_descent(gradient, theta, eta, projection, steps):
    for _ in range(steps):
        theta -= eta * gradient(theta)
        theta = projection(theta)
    return theta
```

4.17.   Newton's Method

$$\theta^{(t+1)} = \theta^{(t)} - [H(\theta^{(t)})]^{-1}\nabla J(\theta^{(t)})$$

Explanation: Newton's method uses second-order information via the Hessian to improve convergence, especially for quadratic cost functions.

Example: For $J(\theta) = \theta^2$, the update uses $H = 2$.

Implementation:

```
def newtons_method(gradient, hessian, theta, steps):
    for _ in range(steps):
        grad = gradient(theta)
        hess = hessian(theta)
        theta -= np.linalg.inv(hess).dot(grad)
    return theta
```

4.18.   Proximal Gradient Method

$$\theta^{(t+1)} = \text{prox}_{\lambda R}(\theta^{(t)} - \eta \nabla J(\theta^{(t)}))$$

Explanation: The proximal gradient method generalizes gradient descent to handle nonsmooth regularization terms such as L1 norm.

Example: For $R(\theta) = \|\theta\|_1$, compute soft thresholding for each parameter.

Implementation:

```
def proximal_gradient(gradient, theta, eta, prox, steps):
    for _ in range(steps):
        theta -= eta * gradient(theta)
        theta = prox(theta)
    return theta
```

4.19.   Proximal Gradient with L1 (ISTA)

$$\theta^{(t+1)} = \text{soft}(\theta^{(t)} - \eta \nabla J(\theta^{(t)}), \lambda\eta)$$

Explanation: Iterative Shrinkage-Thresholding Algorithm (ISTA) applies soft thresholding to update parameters for sparse optimization.

Example: For $J(\theta) = \|\theta - \mathbf{y}\|_2^2 + \lambda\|\theta\|_1$, apply shrinkage to each θ_i.

Implementation:

```
def ista(gradient, theta, eta, lam, steps):
    def soft_threshold(x, lam):
        return np.sign(x) * max(0, abs(x) - lam)
    for _ in range(steps):
        theta -= eta * gradient(theta)
        theta = np.vectorize(soft_threshold)(theta, lam * eta)
    return theta
```

4.20. Penalty Method

$$J_{\text{penalty}}(\theta) = J(\theta) + \frac{1}{\mu}h(\theta)^2$$

Explanation: The penalty method solves constrained optimization problems by penalizing constraint violations in the objective function.

Example: For $h(\theta) = \|\theta\|_2^2 - 1$, penalize deviations from the unit ball constraint.

Implementation:

```python
def penalty_method(loss, theta, penalty, mu):
    return loss(theta) + penalty(theta)**2 / mu
```

4.21.   Augmented Lagrangian Method

$$\mathcal{L}(\theta, \lambda, \mu) = J(\theta) + \lambda h(\theta) + \frac{\mu}{2} h(\theta)^2$$

Explanation: The augmented Lagrangian method combines Lagrangian and penalty approaches to solve constrained optimization problems. It alternates between updating parameters and Lagrange multipliers.

Example: For $J(\theta) = \|\theta\|_2^2$ and $h(\theta) = \|\theta\|_1 - 1$, compute updates for θ, λ, and μ.

Implementation:

```
def augmented_lagrangian(loss, h, theta, lam, mu, steps):
    for _ in range(steps):
        lagrangian = loss(theta) + lam * h(theta) + (mu / 2) * h(theta)**2
        theta -= np.gradient(lagrangian)
        lam += mu * h(theta)
    return theta
```

4.22.   Dual Ascent Method

$$\lambda^{(t+1)} = \lambda^{(t)} + \eta h(\theta^{(t)})$$

Explanation: The dual ascent method optimizes the dual problem of constrained optimization by updating the Lagrange multipliers iteratively.

Example: For $h(\theta) = \|\theta\|_1 - 1$, update λ based on the constraint violation.

Implementation:

```
def dual_ascent(loss, h, theta, lam, eta, steps):
    for _ in range(steps):
        theta -= eta * np.gradient(loss(theta) + lam * h(theta))
        lam += eta * h(theta)
    return theta, lam
```

4.23.   Trust Region Method

$$\theta^{(t+1)} = \arg\min_{\Delta}\{J(\theta) + \nabla J(\theta)^T \Delta + \frac{1}{2}\Delta^T H \Delta \,|\, \|\Delta\| \leq \Delta_{\max}\}$$

Explanation: The trust region method restricts the step size to a region where the quadratic approximation of the cost function is valid, ensuring stability.

Example: For $J(\theta) = \|\theta - \mathbf{y}\|_2^2$, compute steps Δ constrained by $\|\Delta\| \leq \Delta_{\max}$.

Implementation:

```
def trust_region(loss, gradient, hessian, theta, delta_max, steps):
    for _ in range(steps):
        grad = gradient(theta)
        hess = hessian(theta)
        delta = np.linalg.solve(hess, -grad)
        if np.linalg.norm(delta) > delta_max:
            delta *= delta_max / np.linalg.norm(delta)
        theta += delta
    return theta
```

Nabin Katwal

4.24. Barrier Method

$$J_{\text{barrier}}(\theta) = J(\theta) - \frac{1}{\mu} \sum_{i=1}^{m} \ln(-h_i(\theta))$$

Explanation: The barrier method solves constrained optimization by penalizing constraint violations with a logarithmic barrier, keeping updates within the feasible region.

Example: For $h(\theta) = \|\theta\|_1 - 1$, use $-\ln(1 - \|\theta\|_1)$ as the barrier term.

Implementation:

```
def barrier_method(loss, h, theta, mu, steps):
    for _ in range(steps):
        barrier = -np.sum(np.log(-h(theta)))
        theta -= np.gradient(loss(theta) + (1 / mu) * barrier)
        mu *= 0.9
    return theta
```

4.25.   Simulated Annealing

$$P(\Delta E) = \exp\left(-\frac{\Delta E}{T}\right)$$

Explanation: Simulated annealing is a probabilistic optimization algorithm inspired by annealing in metallurgy. It explores the solution space by accepting worse solutions probabilistically to escape local minima.

Example: Minimize $J(\theta) = \theta^2$ with an initial temperature $T = 1$, gradually cooling down.

Implementation:

```
import numpy as np
def simulated_annealing(loss, theta, T, cooling_rate, steps):
    for _ in range(steps):
        new_theta = theta + np.random.uniform(-1, 1, size=theta.shape)
        delta_E = loss(new_theta) - loss(theta)
        if delta_E < 0 or np.exp(-delta_E / T) > np.random.rand():
            theta = new_theta
        T *= cooling_rate
    return theta
```

Chapter 5

REGRESSION

5.1. Linear Regression Hypothesis

$$\hat{y} = \mathbf{X}\boldsymbol{\beta} + \epsilon$$

Explanation: The hypothesis for linear regression assumes that the target variable y is a linear combination of features $\mathbf{X}$, coefficients $\boldsymbol{\beta}$, and an error term ϵ.

Example: For $y = 2x_1 + 3x_2 + \epsilon$, predict y as a linear function of x_1 and x_2.

Implementation:

```
import numpy as np
X = np.array([[1, 2], [3, 4]])
beta = np.array([2, 3])
y_pred = X @ beta
```

5.2. Ordinary Least Squares (OLS)

$$\boldsymbol{\beta} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$$

Explanation: OLS finds the coefficient vector $\boldsymbol{\beta}$ that minimizes the sum of squared residuals between predicted and actual values.

Example: For $\mathbf{X} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $\mathbf{y} = \begin{bmatrix} 5 \\ 11 \end{bmatrix}$, compute $\boldsymbol{\beta}$.

Implementation:

```
beta = np.linalg.inv(X.T @ X) @ X.T @ y
```

5.3. Mean Squared Error (MSE)

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}_i)^2$$

Explanation: MSE quantifies the average squared difference between actual and predicted values. It is a standard loss function in regression.

Example: For $y = [1, 2, 3]$ and $\hat{y} = [1.1, 1.9, 3.2]$, compute the MSE.

Implementation:

```
mse = np.mean((y - y_pred)**2)
```

5.4.   Gradient of the MSE Loss

$$\frac{\partial}{\partial \boldsymbol{\beta}} \text{MSE} = -\frac{2}{n} \mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

Explanation: The gradient of MSE with respect to $\boldsymbol{\beta}$ is used in gradient-based optimization algorithms like gradient descent.

Example: Compute the gradient for $\mathbf{X} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$, $\mathbf{y} = [5, 11]$, and $\boldsymbol{\beta} = [1, 1]$.

Implementation:

```
grad = -2 / len(y) * X.T @ (y - X @ beta)
```

Nabin Katwal

5.5.   Coefficient of Determination (R²)

$$R^2 = 1 - \frac{\sum_{i=1}^{n}(y_i - \hat{y}_i)^2}{\sum_{i=1}^{n}(y_i - \bar{y})^2}$$

Explanation: R^2 measures the proportion of variance in the target variable explained by the model. A value close to 1 indicates a good fit.

Example: For $y = [1, 2, 3]$ and $\hat{y} = [1.1, 1.9, 3.2]$, compute R^2.

Implementation:

```
r2 = 1 - np.sum((y - y_pred)**2) / np.sum((y - np.mean(y))**2)
```

5.6.   Adjusted R²

$$\bar{R}^2 = 1 - \frac{(1 - R^2)(n - 1)}{n - p - 1}$$

Explanation: Adjusted R² accounts for the number of predictors p in the model, penalizing overfitting.

Example: For $R^2 = 0.9$, $n = 100$, and $p = 5$, compute $\bar{R}^2$.

Implementation:

```
adjusted_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)
```

5.7.   Mean Absolute Error (MAE)

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^{n} |y_i - \hat{y}_i|$$

Explanation: MAE measures the average magnitude of prediction errors. It is less sensitive to outliers compared to MSE.

Example: For $y = [1, 2, 3]$ and $\hat{y} = [1.1, 1.9, 3.2]$, compute the MAE.

Implementation:

```
mae = np.mean(np.abs(y - y_pred))
```

Nabin Katwal

5.8.   Weighted Least Squares (WLS)

$$\boldsymbol{\beta} = (\mathbf{X}^T\mathbf{W}\mathbf{X})^{-1}\mathbf{X}^T\mathbf{W}\mathbf{y}$$

Explanation: WLS minimizes the sum of weighted residuals, allowing for heteroscedasticity in the data.

Example: For $\mathbf{W} = \mathrm{diag}([1, 2])$, compute $\boldsymbol{\beta}$.

Implementation:

```
W = np.diag([1, 2])
beta = np.linalg.inv(X.T @ W @ X) @ X.T @ W @ y
```

5.9.   Polynomial Regression Hypothesis

$$\hat{y} = \boldsymbol{\beta}_0 + \boldsymbol{\beta}_1 x + \boldsymbol{\beta}_2 x^2 + \cdots + \boldsymbol{\beta}_n x^n$$

Explanation: Polynomial regression models the relationship between x and y as a polynomial. It generalizes linear regression to non-linear patterns.

Example: Fit $y = 2x + x^2$.

Implementation:

```
from numpy.polynomial.polynomial import Polynomial
poly = Polynomial.fit(X, y, deg=2)
y_pred = poly(X)
```

Nabin Katwal   119

5.10.   Non-Linear Regression

$$\hat{y} = f(\mathbf{X}, \boldsymbol{\beta}) + \epsilon$$

Explanation: Non-linear regression models relationships where the target variable is a non-linear function of the parameters.

Example: Fit $y = ae^{bx}$ using optimization.

Implementation:

```
from scipy.optimize import curve_fit
def model(X, a, b):
    return a * np.exp(b * X)
params, _ = curve_fit(model, X, y)
```

5.11.   Maximum Likelihood Estimation for Regression

$$\hat{\boldsymbol{\beta}} = \arg\max_{\boldsymbol{\beta}} \prod_{i=1}^{n} p(y_i \mid \mathbf{X}_i, \boldsymbol{\beta})$$

Explanation: MLE estimates the parameters that maximize the likelihood of observing the data under a probabilistic model.

Example: Estimate $\boldsymbol{\beta}$ assuming Gaussian noise.

Implementation:

```
from scipy.optimize import minimize
def neg_log_likelihood(beta, X, y):
    residuals = y - X @ beta
    return np.sum(residuals**2)
beta = minimize(neg_log_likelihood, np.zeros(X.shape[1]), args=(X, y)).x
```

5.12.   Empirical Risk Minimization

$$\hat{\boldsymbol{\theta}} = \arg\min_{\boldsymbol{\theta}} \frac{1}{n} \sum_{i=1}^{n} \ell(y_i, f(\mathbf{X}_i, \boldsymbol{\theta}))$$

Explanation: ERM minimizes the average loss over the training data to estimate the model parameters.

Example: Minimize MSE loss for linear regression.

Implementation:

```
def empirical_risk(theta, X, y, loss):
    return np.mean([loss(y[i], np.dot(X[i], theta)) for i in range(len(y))])
```

5.13. Logistic Regression Hypothesis

$$\hat{y} = \sigma(\mathbf{X}\boldsymbol{\beta}), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

Explanation: Logistic regression predicts probabilities for binary classification using the sigmoid function applied to a linear combination of inputs.

Example: For $\mathbf{X} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $\boldsymbol{\beta} = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$, compute $\hat{y}$.

Implementation:

```
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
y_pred = sigmoid(X @ beta)
```

5.14.   Binary Cross-Entropy Loss

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^{n} [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Explanation: Binary cross-entropy measures the dissimilarity between predicted probabilities and true labels in binary classification.

Example: For $y = [1, 0]$ and $\hat{y} = [0.9, 0.1]$, compute the loss.

Implementation:

```
loss = -np.mean(y * np.log(y_pred) + (1 - y) * np.log(1 - y_pred))
```

5.15. Cross-Entropy Loss (Multi-Class)

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^{n} \sum_{j=1}^{k} y_{ij} \log(\hat{y}_{ij})$$

Explanation: Cross-entropy loss generalizes to multi-class classification, comparing one-hot-encoded true labels with predicted probabilities.

Example: For $y = [1, 0, 0]$ and $\hat{y} = [0.8, 0.1, 0.1]$, compute the loss.

Implementation:

```
loss = -np.mean(np.sum(y * np.log(y_pred), axis=1))
```

5.16.   Hinge Loss for SVM

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^{n} \max(0, 1 - y_i \hat{y}_i)$$

Explanation: Hinge loss penalizes predictions that are not at least 1 margin away from the correct classification in support vector machines (SVMs).

Example: For $y = [1, -1]$ and $\hat{y} = [0.8, -0.5]$, compute the loss.

Implementation:

```
loss = np.mean(np.maximum(0, 1 - y * y_pred))
```

5.17.   Lasso Regression Objective

$$\mathcal{L} = \frac{1}{2n}\|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \lambda\|\boldsymbol{\beta}\|_1$$

Explanation: Lasso regression adds an L1 regularization term to the least squares loss, promoting sparsity in the coefficients.

Example: For $\lambda = 0.1$, add $\|\boldsymbol{\beta}\|_1$ as a penalty.

Implementation:

```
loss = 0.5 * np.mean((y - X @ beta)**2) + lam * np.sum(np.abs(beta))
```

5.18.   Ridge Regression Objective

$$\mathcal{L} = \frac{1}{2n}\|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|_2^2 + \lambda\|\boldsymbol{\beta}\|_2^2$$

Explanation: Ridge regression adds an L2 regularization term to reduce overfitting by shrinking coefficients.

Example: For $\lambda = 0.1$, compute the loss with L2 regularization.

Implementation:

```
loss = 0.5 * np.mean((y - X @ beta)**2) + lam * np.sum(beta**2)
```

5.19.   Negative Binomial Regression

$$\hat{y} = \frac{\Gamma(y + \alpha)}{\Gamma(y + 1)\Gamma(\alpha)} \left(\frac{\alpha}{\alpha + \hat{\mu}}\right)^{\alpha} \left(\frac{\hat{\mu}}{\alpha + \hat{\mu}}\right)^{y}$$

Explanation: Negative binomial regression models count data with overdispersion using a generalized linear model.

Example: Fit a model for overdispersed count data.

Implementation:

```
from statsmodels.api import GLM, families
model = GLM(y, X, family=families.NegativeBinomial())
results = model.fit()
```

5.20.   Poisson Regression Model

$$\hat{\mu} = e^{\mathbf{X}\boldsymbol{\beta}}$$

Explanation: Poisson regression models count data using a log link function, assuming the target variable follows a Poisson distribution.

Example: Predict event counts given feature data.

Implementation:

```
from statsmodels.api import GLM, families
model = GLM(y, X, family=families.Poisson())
results = model.fit()
```

Nabin Katwal

5.21. Gamma Regression Objective

$$\mathcal{L} = \frac{1}{\phi} \sum_{i=1}^{n} \left(-\log(\hat{\mu}_i) + \frac{y_i}{\hat{\mu}_i} \right)$$

Explanation: Gamma regression models positive continuous data with a Gamma distribution, often for skewed datasets.

Example: Predict insurance claims amounts.

Implementation:

```
from statsmodels.api import GLM, families
model = GLM(y, X, family=families.Gamma())
results = model.fit()
```

Nabin Katwal

5.22. Probit Regression Model

$$P(y = 1) = \Phi(\mathbf{X}\boldsymbol{\beta})$$

Explanation: Probit regression models binary classification using the cumulative normal distribution function Φ.

Example: Predict binary outcomes using a probit link.

Implementation:

```
from statsmodels.api import GLM, families
model = GLM(y, X, family=families.Binomial(link=families.links.probit()))
results = model.fit()
```

Nabin Katwal

5.23.   Multinomial Logistic Regression

$$P(y = k) = \frac{e^{\mathbf{X}\boldsymbol{\beta}_k}}{\sum_{j=1}^{K} e^{\mathbf{X}\boldsymbol{\beta}_j}}$$

Explanation: Multinomial logistic regression generalizes logistic regression for multi-class classification tasks.

Example: Classify samples into one of $K = 3$ classes.

Implementation:

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(multi_class='multinomial')
model.fit(X, y)
```

5.24. Quantile Regression Loss

$$\mathcal{L} = \sum_{i=1}^{n} \rho_{\tau}(y_i - \hat{y}_i), \quad \rho_{\tau}(e) = \max(\tau e, (1 - \tau)e)$$

Explanation: Quantile regression minimizes the weighted sum of residuals, modeling conditional quantiles of the target variable.

Example: Estimate the 90th percentile of target values.

Implementation:

```
from statsmodels.api import QuantReg
model = QuantReg(y, X)
results = model.fit(q=0.9)
```

Nabin Katwal

5.25.   Huber Loss

$$\mathcal{L} = \sum_{i=1}^{n} \begin{cases} \frac{1}{2}(y_i - \hat{y}_i)^2, & \text{if } |y_i - \hat{y}_i| \leq \delta \\ \delta|y_i - \hat{y}_i| - \frac{1}{2}\delta^2, & \text{otherwise} \end{cases}$$

Explanation: Huber loss combines MSE and MAE, being quadratic for small errors and linear for large errors, robust to outliers.

Example: Fit a regression model robust to outliers with $\delta = 1$.

Implementation:

```
def huber_loss(y, y_pred, delta):
    diff = np.abs(y - y_pred)
    return np.where(diff <= delta, 0.5 * diff**2, delta * diff - 0.5 * delta**2)
```

Nabin Katwal   135

136

Chapter 6

NEURAL NETWORKS

6.1. Perceptron Update Rule

$$\boldsymbol{w}^{(t+1)} = \boldsymbol{w}^{(t)} + \eta(y - \hat{y})\mathbf{x}$$

Explanation: The perceptron update rule adjusts weights based on prediction errors. It is used for binary classification in linearly separable data.

Example: For $\mathbf{x} = [1, 2]$, $y = 1$, $\hat{y} = 0$, and $\eta = 0.1$, update $\boldsymbol{w}$.

Implementation:

```
w += eta * (y - y_pred) * x
```

6.2.   Forward Propagation (Single Layer)

$$\hat{y} = \sigma(\mathbf{X}\boldsymbol{w} + b)$$

Explanation: Forward propagation computes predictions by applying a weight matrix and activation function to input features.

Example: For $\mathbf{X} = [1, 2]$, $\boldsymbol{w} = [0.5, 0.5]$, and $b = 0$, compute $\hat{y}$.

Nabin Katwal

6.3.   Sigmoid Activation

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Explanation: The sigmoid activation maps inputs to $[0, 1]$, commonly used for binary classification.

Example: For $z = 0.5$, compute $\sigma(0.5)$.

Implementation:

```
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

Nabin Katwal

6.4. Tanh Activation

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Explanation: Tanh activation maps inputs to $[-1, 1]$ and is useful for symmetric data.

Example: For $z = 0.5$, compute $\tanh(0.5)$.

Implementation:

```
def tanh(z):
    return np.tanh(z)
```

6.5.   ReLU Activation

$$\text{ReLU}(z) = \max(0, z)$$

Explanation: ReLU introduces non-linearity by zeroing negative values, often used in deep networks.

Example: For $z = -1$, compute ReLU(-1).

Implementation:

```
def relu(z):
    return np.maximum(0, z)
```

6.6. Heaviside Step Activation

$$H(z) = \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases}$$

Explanation: The Heaviside step function outputs binary values for classification tasks.

Example: For $z = -1$, compute $H(-1)$.

Implementation:

```
def heaviside(z):
    return np.where(z >= 0, 1, 0)
```

6.7. Leaky ReLU Activation

$$\text{Leaky ReLU}(z) = \begin{cases} z, & z \geq 0 \\ \alpha z, & z < 0 \end{cases}$$

Explanation: Leaky ReLU allows small gradients for negative inputs, mitigating dead neurons.

Example: For $z = -1$ and $\alpha = 0.01$, compute Leaky ReLU(-1).

Implementation:

```
def leaky_relu(z, alpha=0.01):
    return np.where(z >= 0, z, alpha * z)
```

6.8.   ELU Activation (Exponential Linear Unit)

$$\text{ELU}(z) = \begin{cases} z, & z \geq 0 \\ \alpha(e^z - 1), & z < 0 \end{cases}$$

Explanation: ELU smooths ReLU by providing exponential outputs for negative inputs, improving gradient flow.

Example: For $z = -1$ and $\alpha = 1$, compute ELU(-1).

Implementation:

```
def elu(z, alpha=1):
    return np.where(z >= 0, z, alpha * (np.exp(z) - 1))
```

6.9.   Softmax Function

$$\text{Softmax}(z)_i = \frac{e^{z_i}}{\sum_{j=1}^{n} e^{z_j}}$$

Explanation: Softmax normalizes a vector into a probability distribution over n classes.

Example: For $\mathbf{z} = [1, 2, 3]$, compute Softmax($\mathbf{z}$).

Implementation:

```python
def softmax(z):
    exp_z = np.exp(z - np.max(z))  # Numerical stability
    return exp_z / exp_z.sum(axis=0)
```

Nabin Katwal

6.10. Loss Function for Multi-Class (Cross-Entropy)

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^{n} \sum_{j=1}^{k} y_{ij} \log(\hat{y}_{ij})$$

Explanation: Cross-entropy loss measures the dissimilarity between predicted probabilities and true labels in multi-class classification.

Example: For $y = [1, 0, 0]$ and $\hat{y} = [0.8, 0.1, 0.1]$, compute the loss.

Implementation:

```
loss = -np.mean(np.sum(y * np.log(y_pred), axis=1))
```

6.11.   Gradient Descent for Neural Networks

$$\theta^{(t+1)} = \theta^{(t)} - \eta \frac{\partial \mathcal{L}}{\partial \theta}$$

Explanation: Gradient descent updates the network's weights by minimizing the loss function using gradients.

Example: Update θ for $\mathcal{L} = (y - \hat{y})^2$.

Nabin Katwal

6.12.   Backpropagation (Gradient for Weights)

$$\frac{\partial \mathcal{L}}{\partial w_{ij}} = \delta_j a_i, \quad \delta_j = \frac{\partial \mathcal{L}}{\partial z_j} \sigma'(z_j)$$

Explanation: Backpropagation computes the gradient of the loss function with respect to the weights in a neural network using the chain rule.

Example: Compute gradients for a single-layer neural network.

Implementation:

```
delta = (y_pred - y) * sigmoid_prime(z)
grad_w = np.outer(delta, a)
```

Nabin Katwal   148

6.13.   Mean Squared Error Loss

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}_i)^2$$

Explanation: Mean squared error measures the average squared difference between predictions and actual values, commonly used in regression.

Example: For $y = [1, 2]$ and $\hat{y} = [1.1, 1.8]$, compute the loss.

Implementation:

```
loss = np.mean((y - y_pred)**2)
```

6.14. Binary Cross-Entropy Loss

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^{n} [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Explanation: Binary cross-entropy measures the difference between predicted probabilities and true binary labels.

Example: For $y = [1, 0]$ and $\hat{y} = [0.9, 0.1]$, compute the loss.

Implementation:

```
loss = -np.mean(y * np.log(y_pred) + (1 - y) * np.log(1 - y_pred))
```

Nabin Katwal

6.15.   Batch Normalization

$$\hat{x} = \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}}, \quad y = \gamma\hat{x} + \beta$$

Explanation: Batch normalization normalizes inputs to a layer, reducing internal covariate shift and accelerating training.

Example: Normalize $x = [1, 2, 3]$ with $\gamma = 1, \beta = 0$.

Implementation:

```
mean = np.mean(x)
var = np.var(x)
x_norm = (x - mean) / np.sqrt(var + epsilon)
y = gamma * x_norm + beta
```

Nabin Katwal

6.16. Dropout Regularization

$$\hat{a}_i = \begin{cases} 0, & \text{with probability } p \\ \frac{a_i}{1-p}, & \text{otherwise} \end{cases}$$

Explanation: Dropout randomly sets a fraction p of activations to zero during training to prevent overfitting.

Example: Apply dropout to activations $a = [1, 2, 3]$ with $p = 0.5$.

Implementation:

```
mask = np.random.rand(len(a)) > p
a_dropout = a * mask / (1 - p)
```

6.17.   Gradient of Sigmoid

$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

Explanation: The derivative of the sigmoid function is used in backpropagation to compute gradients efficiently.

Example: For $z = 0.5$, compute $\sigma'(0.5)$.

Implementation:

```python
def sigmoid_prime(z):
    s = sigmoid(z)
    return s * (1 - s)
```

6.18.   RMSProp for Weight Updates

$$s^{(t+1)} = \beta s^{(t)} + (1 - \beta)g^2, \quad w^{(t+1)} = w^{(t)} - \frac{\eta}{\sqrt{s^{(t+1)} + \epsilon}}g$$

Explanation: RMSProp adapts the learning rate for each weight based on the moving average of squared gradients.

Implementation:

```
s = beta * s + (1 - beta) * grad**2
w -= eta / (np.sqrt(s) + epsilon) * grad
```

6.19. Xavier (Glorot) Initialization

$$w \sim \mathcal{U}(-\sqrt{\frac{6}{n_{\mathrm{in}} + n_{\mathrm{out}}}}, \sqrt{\frac{6}{n_{\mathrm{in}} + n_{\mathrm{out}}}})$$

Explanation: Xavier initialization sets weights to maintain variance across layers, improving convergence in deep networks.

Implementation:

```
limit = np.sqrt(6 / (n_in + n_out))
w = np.random.uniform(-limit, limit, size=(n_in, n_out))
```

Nabin Katwal

6.20.   L2 Regularization (Weight Decay)

$$\mathcal{L} = \mathcal{L}_0 + \frac{\lambda}{2}\|\mathbf{w}\|_2^2$$

Explanation: L2 regularization adds a penalty proportional to the square of weights to prevent overfitting.

6.21.   Heaviside vs. Hard Sigmoid

$$\text{Hard Sigmoid}(z) = \max(0, \min(1, 0.2z + 0.5))$$

Explanation: Heaviside is a binary activation function, while Hard Sigmoid approximates sigmoid for efficiency.

Implementation:

```
def hard_sigmoid(z):
    return np.clip(0.2 * z + 0.5, 0, 1)
```

6.22. Swish Activation

$$\text{Swish}(z) = z \cdot \sigma(z)$$

Explanation: Swish is a smooth, non-monotonic activation function that often outperforms ReLU in deep networks.

Implementation:

```
def swish(z):
    return z * sigmoid(z)
```

6.23. Maxout Activation

$$\mathrm{Maxout}(\mathbf{z}) = \max_{i \in [1,k]} z_i$$

Explanation: Maxout selects the maximum value from k linear functions, enabling learnable piecewise linear activations.

Implementation:

```
def maxout(z):
    return np.max(z, axis=0)
```

6.24.   Sparse Categorical Cross-Entropy

$$\mathcal{L} = -\frac{1}{n} \sum_{i=1}^{n} \log(\hat{y}_{i,y_i})$$

Explanation: Sparse categorical cross-entropy simplifies the loss calculation by directly indexing the true class probabilities.

Implementation:

```
loss = -np.mean(np.log(y_pred[range(len(y)), y]))
```

6.25. Cosine Similarity / Cosine Loss

$$\text{Cosine Similarity} = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\| \|\mathbf{v}\|}$$

Explanation: Cosine similarity measures the angle between vectors, commonly used in text and embedding similarity.

Implementation:

```
cos_sim = np.dot(u, v) / (np.linalg.norm(u) * np.linalg.norm(v))
```

162

Chapter 7

CLUSTERING

7.1. Distance Metric (Euclidean)

$$d(\mathbf{u}, \mathbf{v}) = \sqrt{\sum_{i=1}^{n} (u_i - v_i)^2}$$

Explanation: Euclidean distance measures the straight-line distance between two points in n-dimensional space. It is widely used in clustering and nearest-neighbor methods.

Example: For $\mathbf{u} = [1, 2]$ and $\mathbf{v} = [3, 4]$, $d(\mathbf{u}, \mathbf{v}) = \sqrt{(3 - 1)^2 + (4 - 2)^2} = \sqrt{8}$.

Implementation:

```
def euclidean_distance(u, v):
    return np.sqrt(np.sum((u - v)**2))
```

7.2. Manhattan Distance

$$d(\mathbf{u}, \mathbf{v}) = \sum_{i=1}^{n} |u_i - v_i|$$

Explanation: Manhattan distance measures the sum of absolute differences between corresponding components, resembling city block distances.

Example: For $\mathbf{u} = [1, 2]$ and $\mathbf{v} = [3, 4]$, $d(\mathbf{u}, \mathbf{v}) = |3 - 1| + |4 - 2| = 4$.

Implementation:

```
def manhattan_distance(u, v):
    return np.sum(np.abs(u - v))
```

7.3.   Cosine Similarity

$$\text{Cosine Similarity} = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{u}\|\|\mathbf{v}\|}$$

Explanation: Cosine similarity measures the cosine of the angle between two vectors, capturing orientation rather than magnitude.

Example: For $\mathbf{u} = [1, 0]$ and $\mathbf{v} = [0, 1]$, similarity is 0.

Implementation:

```
def cosine_similarity(u, v):
    return np.dot(u, v) / (np.linalg.norm(u) * np.linalg.norm(v))
```

Nabin Katwal

7.4. Jaccard Similarity (Binary Data)

$$\text{Jaccard Similarity} = \frac{|\mathbf{u} \cap \mathbf{v}|}{|\mathbf{u} \cup \mathbf{v}|}$$

Explanation: Jaccard similarity compares the intersection and union of binary data, commonly used in text and set-based similarity.

Example: For $\mathbf{u} = [1, 1, 0]$ and $\mathbf{v} = [1, 0, 1]$, similarity is $\frac{1}{3}$.

Implementation:

```
def jaccard_similarity(u, v):
    return np.sum(np.logical_and(u, v)) / np.sum(np.logical_or(u, v))
```

7.5.   k-Means Objective

$$J = \sum_{i=1}^{k} \sum_{\mathbf{x}_j \in C_i} \|\mathbf{x}_j - \boldsymbol{\mu}_i\|^2$$

Explanation: The k-means objective minimizes the sum of squared distances between data points and their assigned cluster centroids.

Example: For points $[1, 2]$, $[3, 4]$ in cluster C_1 with centroid $[2, 3]$, compute J.

Implementation:

```python
def k_means_objective(X, centroids, labels):
    return np.sum(np.linalg.norm(X - centroids[labels], axis=1)**2)
```

7.6. Centroid Update Rule (k-Means)

$$\boldsymbol{\mu}_i = \frac{1}{|C_i|} \sum_{\mathbf{x}_j \in C_i} \mathbf{x}_j$$

Explanation: The centroid of each cluster is updated as the mean of points assigned to it.

Example: For cluster $C_1 = \{[1,2], [3,4]\}$, compute $\boldsymbol{\mu}_1 = [2,3]$.

Implementation:

```
def update_centroids(X, labels, k):
    return np.array([X[labels == i].mean(axis=0) for i in range(k)])
```

Nabin Katwal

7.7. Elbow Method for Optimal k

$$J(k) = \sum_{i=1}^{k} \sum_{\mathbf{x}_j \in C_i} \|\mathbf{x}_j - \boldsymbol{\mu}_i\|^2$$

Explanation: The elbow method finds the optimal number of clusters k by identifying the "elbow" in the plot of $J(k)$ versus k.

Implementation:

```
def elbow_method(X, max_k):
    distortions = []
    for k in range(1, max_k + 1):
        kmeans = KMeans(n_clusters=k).fit(X)
        distortions.append(kmeans.inertia_)
    return distortions
```

Nabin Katwal

7.8. k-Medoids Objective

$$J = \sum_{i=1}^{k} \sum_{\mathbf{x}_j \in C_i} d(\mathbf{x}_j, \mathbf{m}_i)$$

Explanation: k-Medoids minimizes the sum of distances between data points and their cluster medoids, robust to outliers.

Example: Replace centroids with medoids for robust clustering.

Implementation:

```
def k_medoids_objective(X, medoids, labels):
    return np.sum([np.sum(np.linalg.norm(X[labels == i]
    - medoids[i], axis=1)) for i in range(len(medoids))])
```

7.9.   Fuzzy c-Means Objective

$$J = \sum_{i=1}^{c} \sum_{j=1}^{n} u_{ij}^{m} \|\mathbf{x}_j - \mathbf{c}_i\|^2$$

Explanation: Fuzzy c-means assigns membership values u_{ij} to each data point for each cluster, allowing soft clustering.

Implementation:

```
def fuzzy_c_means_objective(X, centroids, memberships, m):
    return np.sum(memberships**m * np.linalg.norm(X[:, None]
    - centroids, axis=2)**2)
```

7.10.   Silhouette Score

$$S = \frac{b - a}{\max(a, b)}, \quad a = \text{intra-cluster distance}, b = \text{nearest-cluster distance}$$

Explanation: Silhouette score evaluates the quality of clustering by comparing intra-cluster and nearest-cluster distances.

Implementation:

```
from sklearn.metrics import silhouette_score
score = silhouette_score(X, labels)
```

7.11.   Hierarchical Clustering Dendrogram

$$d(C_1, C_2) = \min_{x \in C_1, y \in C_2} \|x - y\|$$

Explanation: A dendrogram visually represents the hierarchical clustering process, showing cluster merges.

Implementation:

```
from scipy.cluster.hierarchy import dendrogram, linkage
Z = linkage(X, method='ward')
dendrogram(Z)
```

Nabin Katwal

7.12. Ward's Linkage

$$d(C_1, C_2) = \frac{|C_1||C_2|}{|C_1| + |C_2|} \|\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2\|^2$$

Explanation: Ward's linkage minimizes the variance increase when merging clusters, resulting in compact clusters.

Implementation:

```
from scipy.cluster.hierarchy import linkage
Z = linkage(X, method='ward')
```

7.13.   Single vs. Complete Linkage

$$d_{\text{single}}(C_1, C_2) = \min_{x \in C_1, y \in C_2} \|x - y\|, \quad d_{\text{complete}}(C_1, C_2) = \max_{x \in C_1, y \in C_2} \|x - y\|$$

Explanation: Single linkage merges clusters based on the smallest distance between points, while complete linkage uses the largest distance. They influence the shape of hierarchical clustering.

Implementation:

```
from scipy.cluster.hierarchy import import linkage
Z_single = linkage(X, method='single')
Z_complete = linkage(X, method='complete')
```

7.14.   Average Linkage

$$d_{\text{average}}(C_1, C_2) = \frac{1}{|C_1||C_2|} \sum_{x \in C_1} \sum_{y \in C_2} \|x - y\|$$

Explanation: Average linkage computes the average distance between all pairs of points in two clusters, balancing the extremes of single and complete linkage.

Implementation:

```
Z_average = linkage(X, method='average')
```

7.15.   Minimum Spanning Tree Criterion

$$\text{MST weight} = \sum_{(u,v)\in E} w(u,v), \quad w(u,v) = \|u - v\|$$

Explanation: The minimum spanning tree (MST) connects all points with the minimum total edge weight, often used in clustering to detect dense regions.

Implementation:

```
from scipy.sparse.csgraph import minimum_spanning_tree
mst = minimum_spanning_tree(distance_matrix(X))
```

7.16.   DBSCAN Core Point Condition

$$|\text{Neighbors}(\mathbf{x})| \geq \text{MinPts}, \quad \text{where Neighbors}(\mathbf{x}) = \{\mathbf{y} : \|\mathbf{x} - \mathbf{y}\| \leq \epsilon\}$$

Explanation: A core point in DBSCAN must have at least MinPts neighbors within a distance ϵ.

Implementation:

```
core_condition = len(neighbors) >= MinPts
```

7.17.   DBSCAN Density Condition

Density-connected: $\exists$ a chain of points $\mathbf{x}_1, \mathbf{x}_2, \ldots, \mathbf{x}_n$ such that $\|\mathbf{x}_i - \mathbf{x}_{i+1}\| \leq \epsilon$

Explanation: DBSCAN forms clusters by connecting points that are density-reachable through chains of neighbors.

Implementation:

```
from sklearn.cluster import DBSCAN
dbscan = DBSCAN(eps=epsilon, min_samples=MinPts).fit(X)
```

Nabin Katwal

7.18.   Cohesion Metric

$$\text{Cohesion} = \sum_{i=1}^{k} \sum_{\mathbf{x}_j \in C_i} \|\mathbf{x}_j - \boldsymbol{\mu}_i\|^2$$

Explanation: Cohesion measures the compactness of clusters, where smaller values indicate tighter clusters.

Implementation:

```
cohesion = sum(np.linalg.norm(X[labels == i]
- centroids[i], axis=1).sum() for i in range(k))
```

7.19.   Separation Metric

$$\text{Separation} = \sum_{i=1}^{k} \sum_{j=i+1}^{k} \|\boldsymbol{\mu}_i - \boldsymbol{\mu}_j\|^2$$

Explanation: Separation measures the distance between cluster centroids, where larger values indicate well-separated clusters.

Implementation:

```
separation = sum(np.linalg.norm(centroids[i]
- centroids[j])**2 for i in range(k) for j in range(i+1, k))
```

7.20.   Soft Clustering Membership

$$u_{ij} = \frac{\|\mathbf{x}_j - \mathbf{c}_i\|^{-2/(m-1)}}{\sum_{k=1}^{c} \|\mathbf{x}_j - \mathbf{c}_k\|^{-2/(m-1)}}$$

Explanation: Soft clustering assigns membership values u_{ij} to each point for each cluster, indicating the degree of belonging.

Implementation:

```
memberships = 1 / (distances**(2/(m-1)) / distances.sum(axis=1, keepdims=True))
```

7.21.   Entropy for Clustering Evaluation

$$H = -\sum_{i=1}^{k}\sum_{j=1}^{n} P_{ij} \log P_{ij}$$

Explanation: Entropy measures the uncertainty in clustering assignments, where lower values indicate clearer clustering.

Implementation:

```
entropy = -np.sum(P * np.log(P))
```

Nabin Katwal

7.22.   Mutual Information for Clustering

$$I(U,V) = \sum_{i=1}^{|U|} \sum_{j=1}^{|V|} P_{ij} \log \frac{P_{ij}}{P_i P_j}$$

Explanation: Mutual information measures the shared information between true and predicted clusters.

Implementation:

```
from sklearn.metrics import mutual_info_score
mi = mutual_info_score(true_labels, predicted_labels)
```

7.23.   F-Measure for Clustering

$$F = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Explanation: The F-measure evaluates clustering performance by balancing precision and recall.

Implementation:

```
from sklearn.metrics import f1_score
f_measure = f1_score(true_labels, predicted_labels, average='weighted')
```

7.24.   Adjusted Rand Index (ARI)

$$\text{ARI} = \frac{\text{Index} - \text{Expected Index}}{\text{Max Index} - \text{Expected Index}}$$

Explanation: ARI adjusts the Rand Index for chance, measuring clustering similarity.

Implementation:

```
from sklearn.metrics import adjusted_rand_score
ari = adjusted_rand_score(true_labels, predicted_labels)
```

7.25. Normalized Mutual Information (NMI)

$$\text{NMI} = \frac{2I(U,V)}{H(U) + H(V)}$$

Explanation: NMI normalizes mutual information to compare clustering solutions of different sizes.

Implementation:

```
from sklearn.metrics import normalized_mutual_info_score
nmi = normalized_mutual_info_score(true_labels, predicted_labels)
```

Chapter 8

DIMENSIONALITY REDUCTION

8.1.   Principal Component Analysis (PCA) Objective

$$\text{Maximize: } \text{Var}(\mathbf{z}) = \mathbf{w}^T \mathbf{S} \mathbf{w}, \quad \text{subject to } \|\mathbf{w}\|_2 = 1$$

Explanation: PCA seeks directions (principal components) that maximize the variance of projected data while being orthogonal to each other.

Implementation:

```
from sklearn.decomposition import PCA
pca = PCA(n_components=k).fit(X)
```

8.2. Covariance Matrix for PCA

$$\mathbf{S} = \frac{1}{n-1}(\mathbf{X} - \bar{\mathbf{X}})^T(\mathbf{X} - \bar{\mathbf{X}})$$

Explanation: The covariance matrix captures pairwise feature dependencies and is central to PCA.

Implementation:

```
mean_X = np.mean(X, axis=0)
cov_matrix = np.cov(X - mean_X, rowvar=False)
```

8.3.   Eigen Decomposition for PCA

$$\mathbf{Sw} = \lambda \mathbf{w}$$

Explanation: PCA uses eigen decomposition of the covariance matrix to find eigen-values (variances) and eigenvectors (principal components).

Implementation:

```
eigenvalues, eigenvectors = np.linalg.eig(cov_matrix)
```

8.4. SVD (Singular Value Decomposition)

$$\mathbf{X} = \mathbf{U}\boldsymbol{\Sigma}\mathbf{V}^{T}$$

Explanation: SVD factorizes a matrix into orthogonal components, enabling dimensionality reduction by truncating $\boldsymbol{\Sigma}$.

Implementation:

```
U, S, Vt = np.linalg.svd(X, full_matrices=False)
```

8.5. Reconstruction Error for PCA

$$\text{Error} = \|\mathbf{X} - \hat{\mathbf{X}}\|_F^2, \quad \hat{\mathbf{X}} = \mathbf{Z}\mathbf{W}^T + \bar{\mathbf{X}}$$

Explanation: Reconstruction error quantifies the information loss when reducing dimensionality with PCA.

Implementation:

```
X_hat = Z @ W.T + mean_X
reconstruction_error = np.linalg.norm(X - X_hat, 'fro')**2
```

8.6. Explained Variance Ratio

$$\text{Explained Variance Ratio} = \frac{\lambda_i}{\sum_{j=1}^{n} \lambda_j}$$

Explanation: The explained variance ratio quantifies the proportion of variance captured by each principal component.

Implementation:

```
explained_variance_ratio = eigenvalues / np.sum(eigenvalues)
```

8.7. Cumulative Explained Variance

$$\text{Cumulative Explained Variance} = \sum_{i=1}^{k} \frac{\lambda_i}{\sum_{j=1}^{n} \lambda_j}$$

Explanation: Cumulative explained variance evaluates the total variance captured by the first k principal components.

Implementation:

```
cumulative_explained_variance = np.cumsum(explained_variance_ratio)
```

8.8. Random Projection

$$\mathbf{X}_{\text{proj}} = \mathbf{X}\mathbf{R}, \quad \mathbf{R} \sim \mathcal{N}(0, 1)$$

Explanation: Random projection reduces dimensionality by projecting data onto a lower-dimensional random matrix while approximately preserving distances.

Implementation:

```
from sklearn.random_projection import GaussianRandomProjection
rp = GaussianRandomProjection(n_components=k).fit_transform(X)
```

8.9.   Isomap Distance Matrix

$$d_{ij} = \text{Shortest Path Distance on } \mathcal{G}, \quad \mathcal{G} = (\mathbf{X}, \epsilon\text{-Neighborhoods})$$

Explanation: Isomap computes geodesic distances in a graph of nearest neighbors to preserve non-linear structures in the data.

Implementation:

```
from sklearn.manifold import Isomap
isomap = Isomap(n_neighbors=k).fit_transform(X)
```

Nabin Katwal

8.10. MDS Stress Function

$$\text{Stress} = \sum_{i<j} \left(d_{ij} - \hat{d}_{ij} \right)^2$$

Explanation: The stress function measures the discrepancy between original and embedded distances in Multidimensional Scaling (MDS).

Implementation:

```
from sklearn.manifold import MDS
mds = MDS(n_components=2).fit_transform(X)
```

8.11.   Multidimensional Scaling (MDS)

$$\mathbf{X}_{\text{MDS}} = \arg\min_{\mathbf{Y}} \text{Stress}(\mathbf{Y})$$

Explanation: MDS embeds data into a lower-dimensional space while preserving pairwise distances as much as possible.

Implementation:

```
from sklearn.manifold import MDS
mds = MDS(n_components=k).fit_transform(X)
```

8.12.   NMF (Non-Negative Matrix Factorization)

$$X \approx WH, \quad W \geq 0, \, H \geq 0$$

Explanation: NMF factorizes a non-negative matrix into two lower-rank non-negative matrices, often used in topic modeling and image processing.

Implementation:

```
from sklearn.decomposition import NMF
nmf = NMF(n_components=k).fit_transform(X)
```

8.13.   ICA (Independent Component Analysis) Objective

$$\text{Maximize:} \quad \sum_{i=1}^{n} \log p(s_i), \quad \text{where } \mathbf{s} = \mathbf{WX}$$

Explanation: ICA separates mixed signals into statistically independent components by maximizing non-Gaussianity.

Implementation:

```
from sklearn.decomposition import FastICA
ica = FastICA(n_components=k).fit_transform(X)
```

8.14. Factor Analysis Model

$$\mathbf{X} = \mathbf{Z}\boldsymbol{\Lambda} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(0, \Psi)$$

Explanation: Factor analysis models observed variables as linear combinations of latent factors plus noise.

Implementation:

```
from sklearn.decomposition import FactorAnalysis
fa = FactorAnalysis(n_components=k).fit_transform(X)
```

8.15.   Kernel PCA Transformation

$$\mathbf{K} = \phi(\mathbf{X})\phi(\mathbf{X})^T, \quad \text{Eigen Decomposition: } \mathbf{K}\alpha = \lambda\alpha$$

Explanation: Kernel PCA applies PCA in a high-dimensional feature space defined by a kernel function.

Implementation:

```
from sklearn.decomposition import KernelPCA
kpca = KernelPCA(kernel='rbf', n_components=k).fit_transform(X)
```

Nabin Katwal   203

8.16.   LDA (Fisher's Criterion)

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

Explanation: LDA finds a projection that maximizes class separation by optimizing the ratio of between-class to within-class variance.

Implementation:

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
lda = LinearDiscriminantAnalysis(n_components=k).fit_transform(X, y)
```

8.17. Robust PCA (RPCA)

$$X = L + S, \quad \|L\|_* + \lambda \|S\|_1$$

Explanation: RPCA decomposes a matrix into a low-rank component (L) and a sparse component (S).

Implementation:

```
from r_pca import R_pca
rpca = R_pca(X)
L, S = rpca.fit()
```

8.18.   Hessian LLE

$$\text{Minimize: } \|\mathbf{WX} - \mathbf{X}\|_2^2, \quad \text{subject to local Hessian alignment}$$

Explanation: Hessian LLE preserves local geometric structures while optimizing a low-dimensional embedding.

Implementation:

```
from sklearn.manifold import LocallyLinearEmbedding
hessian_lle = LocallyLinearEmbedding(n_neighbors=k,
method='hessian').fit_transform(X)
```

8.19.   Laplacian Eigenmaps Objective

$$\text{Minimize: } \sum_{i,j} w_{ij}\|\mathbf{y}_i - \mathbf{y}_j\|^2, \quad \mathbf{W} = \text{Graph Weights}$$

Explanation: Laplacian Eigenmaps embeds data while preserving local neighborhood information based on a graph structure.

Implementation:

```
from sklearn.manifold import SpectralEmbedding
laplacian = SpectralEmbedding(n_components=k).fit_transform(X)
```

8.20.   Autoencoder Reconstruction

$$\hat{\mathbf{X}} = \text{Decoder}(\text{Encoder}(\mathbf{X}))$$

Explanation: Autoencoders minimize reconstruction error by compressing data into a latent representation and reconstructing it.

Implementation:

```
from keras.models import Model
encoded = encoder(X)
decoded = decoder(encoded)
```

8.21.   Autoencoder Latent Representation

$$\mathbf{Z} = \text{Encoder}(\mathbf{X})$$

Explanation: The latent representation ($\mathbf{Z}$) compresses input data into a lower-dimensional space for downstream tasks.

Implementation:

```
latent_representation = encoder.predict(X)
```

8.22.   Sparse PCA Objective

$$\text{Maximize: } \|\mathbf{X}\mathbf{W}\|_2^2, \quad \text{subject to sparsity constraints on } \mathbf{W}$$

Explanation: Sparse PCA introduces sparsity in the principal components to improve interpretability.

Implementation:

```
from sklearn.decomposition import SparsePCA
spca = SparsePCA(n_components=k).fit_transform(X)
```

8.23.   t-SNE Objective

$$\text{Minimize: } KL(P||Q) = \sum_{i \neq j} P_{ij} \log \frac{P_{ij}}{Q_{ij}}$$

Explanation: t-SNE minimizes the Kullback-Leibler divergence between high-dimensional and low-dimensional distributions.

Implementation:

```
from sklearn.manifold import TSNE
tsne = TSNE(n_components=k).fit_transform(X)
```

Nabin Katwal

8.24.   Gradient of t-SNE

$$\frac{\partial KL}{\partial y_i} = 4 \sum_{j} (P_{ij} - Q_{ij})(y_i - y_j)Q_{ij}$$

Explanation: The gradient of the t-SNE objective updates low-dimensional embeddings to align distributions.

8.25. UMAP (Uniform Manifold Approximation and Projection)

$$\text{Optimize: } \sum_{i,j} w_{ij}\|y_i - y_j\|^2 - \lambda \sum_{k,l} w_{kl} \log(\|y_k - y_l\|)$$

Explanation: UMAP preserves local and global structures by optimizing a balance between distances and densities.

Implementation:

```
import umap
umap_embedding = umap.UMAP(n_components=k).fit_transform(X)
```

214

Chapter 9

PROBABILITY DISTRIBUTIONS

9.1. Bernoulli Distribution

$$P(X = x) = p^x(1 - p)^{1-x}, \quad x \in \{0, 1\}, 0 \le p \le 1$$

Explanation: The Bernoulli distribution models a single binary event, with success probability p.

Example: For $p = 0.7$, $P(X = 1) = 0.7$, $P(X = 0) = 0.3$.

Implementation:

```
from scipy.stats import bernoulli
prob = bernoulli.pmf(k=1, p=0.7)
```

9.2.   Binomial Distribution

$$P(X = k) = \binom{n}{k} p^k (1-p)^{n-k}, \quad k \in \{0, 1, \ldots, n\}$$

Explanation: The Binomial distribution models the number of successes in n independent Bernoulli trials.

Example: For $n = 5$ and $p = 0.5$, $P(X = 3) = \binom{5}{3}(0.5)^3(0.5)^2 = 0.3125$.

Implementation:

```
from scipy.stats import binom
prob = binom.pmf(k=3, n=5, p=0.5)
```

9.3. Poisson Distribution

$$P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k \in \{0, 1, 2, \ldots\}$$

Explanation: The Poisson distribution models the number of events in a fixed interval, with a mean rate λ.

Example: For $\lambda = 3$, $P(X = 2) = \frac{3^2 e^{-3}}{2!} = 0.224$.

Implementation:

```
from scipy.stats import poisson
prob = poisson.pmf(k=2, mu=3)
```

9.4. Uniform Distribution (Continuous)

$$f(x) = \frac{1}{b-a}, \quad x \in [a,b]$$

Explanation: The continuous uniform distribution assigns equal probability density to all points in $[a,b]$.

Example: For $a = 0$, $b = 2$, $f(1) = \frac{1}{2}$.

Implementation:

```
from scipy.stats import uniform
prob = uniform.pdf(x=1, loc=0, scale=2)
```

9.5.   Discrete Uniform Distribution

$$P(X = x) = \frac{1}{n}, \quad x \in \{1, 2, \ldots, n\}$$

Explanation: The discrete uniform distribution assigns equal probability to n discrete outcomes.

Example: For $n = 6$, $P(X = 3) = \frac{1}{6}$.

Implementation:

```
from scipy.stats import randint
prob = randint.pmf(k=3, low=1, high=7)
```

Nabin Katwal

9.6. Normal (Gaussian) Distribution

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}}e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Explanation: The normal distribution models data with a symmetric bell shape, defined by mean μ and standard deviation σ.

Example: For $\mu = 0$, $\sigma = 1$, $f(0) = \frac{1}{\sqrt{2\pi}} \approx 0.398$.

Implementation:

```
from scipy.stats import norm
prob = norm.pdf(x=0, loc=0, scale=1)
```

9.7. Exponential Distribution

$$f(x) = \lambda e^{-\lambda x}, \quad x \geq 0$$

Explanation: The exponential distribution models the time between events in a Poisson process.

Example: For $\lambda = 2$, $f(1) = 2e^{-2} \approx 0.271$.

Implementation:

```
from scipy.stats import expon
prob = expon.pdf(x=1, scale=1/2)
```

9.8. Geometric Distribution

$$P(X = k) = (1 - p)^{k-1}p, \quad k \in \{1, 2, \ldots\}$$

Explanation: The geometric distribution models the number of trials until the first success in repeated Bernoulli trials.

Example: For $p = 0.5$, $P(X = 3) = (0.5)^2(0.5) = 0.125$.

Implementation:

```
from scipy.stats import geom
prob = geom.pmf(k=3, p=0.5)
```

9.9. Hypergeometric Distribution

$$P(X = k) = \frac{\binom{K}{k}\binom{N-K}{n-k}}{\binom{N}{n}}$$

Explanation: The hypergeometric distribution models successes in n draws without replacement from a population of N with K successes.

Example: For $N = 20$, $K = 7$, $n = 5$, $P(X = 3)$.

Implementation:

```
from scipy.stats import hypergeom
prob = hypergeom.pmf(k=3, M=20, n=5, N=7)
```

9.10. Beta Distribution

$$f(x) = \frac{x^{\alpha-1}(1-x)^{\beta-1}}{B(\alpha, \beta)}, \quad x \in [0, 1]$$

Explanation: The Beta distribution models probabilities as a function of parameters α and β.

Example: For $\alpha = 2$, $\beta = 3$, compute $f(0.5)$.

Implementation:

```
from scipy.stats import beta
prob = beta.pdf(x=0.5, a=2, b=3)
```

9.11.   Gamma Distribution

$$f(x) = \frac{\beta^\alpha x^{\alpha-1} e^{-\beta x}}{\Gamma(\alpha)}, \quad x > 0$$

Explanation: The Gamma distribution generalizes the exponential distribution, often used for waiting times.

Example: For $\alpha = 2$, $\beta = 1$, compute $f(1)$.

Implementation:

```
from scipy.stats import gamma
prob = gamma.pdf(x=1, a=2, scale=1/1)
```

9.12. Multinomial Distribution

$$P(X_1 = k_1, \ldots, X_k = k_k) = \frac{n!}{k_1! \cdots k_k!} p_1^{k_1} \cdots p_k^{k_k}$$

Explanation: The multinomial distribution generalizes the binomial distribution for multiple categories.

Example: For $n = 3$, $\mathbf{p} = [0.2, 0.5, 0.3]$, and $\mathbf{k} = [1, 1, 1]$.

Implementation:

```python
from scipy.stats import multinomial
prob = multinomial.pmf(x=[1, 1, 1], n=3, p=[0.2, 0.5, 0.3])
```

9.13. Chi-Square Distribution

$$f(x) = \frac{x^{k/2-1}e^{-x/2}}{2^{k/2}\Gamma(k/2)}, \quad x > 0$$

Explanation: The chi-square distribution models the sum of squares of k independent standard normal variables, commonly used in hypothesis testing.

Example: For $k = 3$, compute $f(2)$.

Implementation:

```
from scipy.stats import chi2
prob = chi2.pdf(x=2, df=3)
```

9.14.   Student's t-Distribution

$$f(x) = \frac{\Gamma((\nu + 1)/2)}{\sqrt{\nu\pi}\Gamma(\nu/2)} \left(1 + \frac{x^2}{\nu}\right)^{-(\nu+1)/2}$$

Explanation: The Student's t-distribution is used for estimating population parameters when the sample size is small.

Example: For $\nu = 5$, compute $f(1)$.

Implementation:

```
from scipy.stats import t
prob = t.pdf(x=1, df=5)
```

9.15.   F-Distribution

$$f(x) = \frac{\sqrt{\left(\frac{d_1 x}{d_2}\right)^{d_1} \left(1 + \frac{d_1 x}{d_2}\right)^{-(d_1 + d_2)/2}}}{xB(d_1/2, d_2/2)}, \quad x > 0$$

Explanation: The F-distribution models the ratio of variances and is commonly used in ANOVA tests.

Implementation:

```
from scipy.stats import f
prob = f.pdf(x=2, dfn=5, dfd=10)
```

Nabin Katwal

9.16. Laplace Distribution

$$f(x) = \frac{1}{2b} e^{-\frac{|x-\mu|}{b}}$$

Explanation: The Laplace distribution, also known as the double exponential distribution, is used for modeling differences in data.

Implementation:

```
from scipy.stats import laplace
prob = laplace.pdf(x=0, loc=0, scale=1)
```

9.17.   Rayleigh Distribution

$$f(x) = \frac{x}{\sigma^2} e^{-x^2/(2\sigma^2)}, \quad x \geq 0$$

Explanation: The Rayleigh distribution models the magnitude of a two-dimensional vector with independent normal components.

Implementation:

```
from scipy.stats import rayleigh
prob = rayleigh.pdf(x=2, scale=1)
```

9.18.   Triangular Distribution

$$f(x) = \begin{cases} \frac{2(x-a)}{(b-a)(c-a)}, & a \leq x < c \\ \\ \frac{2(b-x)}{(b-a)(b-c)}, & c \leq x \leq b \end{cases}$$

Explanation: The triangular distribution models data with a known minimum, maximum, and mode.

Implementation:

```
from scipy.stats import triang
prob = triang.pdf(x=0.5, c=0.5, loc=0, scale=1)
```

9.19. Log-Normal Distribution

$$f(x) = \frac{1}{x\sigma\sqrt{2\pi}} e^{-\frac{(\ln x - \mu)^2}{2\sigma^2}}, \quad x > 0$$

Explanation: The log-normal distribution models data whose logarithm follows a normal distribution.

Implementation:

```
from scipy.stats import lognorm
prob = lognorm.pdf(x=2, s=1, scale=np.exp(0))
```

9.20. Arcsine Distribution

$$f(x) = \frac{1}{\pi\sqrt{x(1-x)}}, \quad x \in (0,1)$$

Explanation: The arcsine distribution models probabilities with endpoints more likely than the middle.

Implementation:

```
from scipy.stats import arcsine
prob = arcsine.pdf(x=0.5)
```

9.21.   Beta-Binomial Distribution

$$P(X = k) = \binom{n}{k} \frac{B(k + \alpha, n - k + \beta)}{B(\alpha, \beta)}$$

Explanation: The beta-binomial distribution models overdispersed binomial outcomes using a Beta prior.

Implementation:

```
from scipy.stats import betabinom
prob = betabinom.pmf(k=2, n=5, a=2, b=3)
```

9.22.   Cauchy Distribution

$$f(x) = \frac{1}{\pi\gamma\left[1 + \left(\frac{x-x_0}{\gamma}\right)^2\right]}$$

Explanation: The Cauchy distribution models data with heavy tails, often used in robust statistics.

Implementation:

```
from scipy.stats import cauchy
prob = cauchy.pdf(x=0, loc=0, scale=1)
```

9.23.   Weibull Distribution

$$f(x) = \frac{k}{\lambda} \left(\frac{x}{\lambda}\right)^{k-1} e^{-(x/\lambda)^k}, \quad x \geq 0$$

Explanation: The Weibull distribution is used for reliability analysis and modeling lifetimes.

Implementation:

```
from scipy.stats import weibull_min
prob = weibull_min.pdf(x=2, c=1.5, scale=1)
```

9.24.   Pareto Distribution

$$f(x) = \frac{\alpha x_m^{\alpha}}{x^{\alpha+1}}, \quad x \geq x_m$$

Explanation: The Pareto distribution models wealth distribution and heavy-tailed phenomena.

Implementation:

```
from scipy.stats import pareto
prob = pareto.pdf(x=2, b=1)
```

9.25.   Log-Cauchy Distribution

$$f(x) = \frac{1}{x\pi\gamma \left[1 + \left(\frac{\ln x - x_0}{\gamma}\right)^2\right]}, \quad x > 0$$

Explanation: The log-Cauchy distribution is the logarithmic transform of the Cauchy distribution, with heavy tails.

Chapter 10

REINFORCEMENT LEARNING

10.1.   Reward Function

$$R(s, a) = \mathbb{E}[\text{Reward} \mid s, a]$$

Explanation: The reward function provides the immediate reward received after taking action a in state s, guiding the agent's behavior.

Implementation:

```
def reward_function(state, action):
    # Example reward calculation
    return rewards[state, action]
```

10.2. Discounted Return

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}, \quad 0 \le \gamma < 1$$

Explanation: The discounted return accumulates rewards over time, weighting future rewards by the discount factor γ.

Implementation:

```python
def discounted_return(rewards, gamma):
    G = 0
    for t, r in enumerate(rewards):
        G += (gamma**t) * r
    return G
```

10.3.   Bellman Equation (State-Value Function)

$$V(s) = \mathbb{E}_\pi[R(s, a) + \gamma V(s')]$$

Explanation: The Bellman equation relates the value of a state to the expected return from it under a policy π.

Implementation:

```
def bellman_state_value(s, rewards, transition_prob, gamma, V):
    return np.sum(transition_prob[s] * (rewards[s] + gamma * V))
```

Nabin Katwal

10.4. Bellman Equation (Action-Value Function)

$$Q(s, a) = \mathbb{E}[R(s, a) + \gamma V(s')]$$

Explanation: The Bellman equation for the action-value function expresses the value of taking action a in state s and following the policy afterward.

Implementation:

```
def bellman_action_value(s, a, rewards, transition_prob, gamma, V):
    return rewards[s, a] + gamma * np.sum(transition_prob[s, a] * V)
```

10.5.   Temporal Difference (TD) Update

$$V(s_t) \leftarrow V(s_t) + \alpha \left[R_{t+1} + \gamma V(s_{t+1}) - V(s_t) \right]$$

Explanation: The TD update improves the value estimate of a state by using the difference between predicted and actual returns.

Implementation:

```
def td_update(V, state, reward, next_state, alpha, gamma):
    V[state] += alpha * (reward + gamma * V[next_state] - V[state])
```

10.6.   Monte Carlo Policy Evaluation

$$V(s) \leftarrow \mathbb{E}[G_t \mid s_t = s]$$

Explanation: Monte Carlo evaluation updates the value of a state by averaging returns from multiple episodes starting from that state.

Implementation:

```
def monte_carlo_evaluation(V, state_returns, state_counts):
    for state, returns in state_returns.items():
        V[state] = np.mean(returns)
```

10.7.   Policy Improvement

$$\pi'(s) = \arg\max_a Q(s, a)$$

Explanation: Policy improvement updates the policy by choosing the action that maximizes the action-value function.

Implementation:

```
def policy_improvement(Q):
    return np.argmax(Q, axis=1)
```

Nabin Katwal

10.8. Q-Learning Update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right]$$

Explanation: Q-learning is an off-policy algorithm that updates action-value estimates using the maximum future Q-value.

Implementation:

```
def q_learning_update(Q, state, action, reward, next_state, alpha, gamma):
    Q[state, action] += alpha * (reward + gamma * np.max(Q[next_state])
    - Q[state, action])
```

10.9.   SARSA Update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t) \right]$$

Explanation: SARSA is an on-policy algorithm that updates Q-values based on the action actually taken under the current policy.

Implementation:

```
def sarsa_update(Q, state, action, reward,
next_state, next_action, alpha, gamma):
    Q[state, action] += alpha * (reward + gamma * Q[next_state, next_action]
    - Q[state, action])
```

10.10. Value Iteration Update

$$V(s) \leftarrow \max_a \left[R(s,a) + \gamma \sum_{s'} P(s' \mid s,a)V(s') \right]$$

Explanation: Value iteration iteratively updates state values by finding the optimal action at each step.

Implementation:

```
def value_iteration(V, rewards, transition_prob, gamma):
    for s in range(len(V)):
        V[s] = max(np.sum(transition_prob[s, a] * (rewards[s, a]
        + gamma * V)) for a in range(num_actions))
```

10.11.   Actor–Critic Policy Update

$$\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(a_t \mid s_t)\delta_t, \quad \delta_t = R_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

Explanation: The actor updates the policy using the advantage, while the critic updates the value function to estimate the advantage.

Implementation:

```
def actor_critic_update(actor, critic, state, action, reward, next_state,
alpha, gamma):
    delta = reward + gamma * critic[next_state] - critic[state]
    actor.update(state, action, alpha * delta)
    critic[state] += alpha * delta
```

10.12.   Deterministic Policy Gradient

$$\nabla J(\theta) = \mathbb{E}_{s \sim \rho^\pi} \left[\nabla_a Q(s, a) \nabla_\theta \pi_\theta(s) \right]$$

Explanation: Deterministic policy gradients update the policy directly in a continuous action space using gradients of the Q-function.

Implementation:

```
def deterministic_policy_gradient(policy, q_function, state, alpha):
    action = policy(state)
    grad_q = q_function.gradient(state, action)
    grad_pi = policy.gradient(state)
    policy.update(state, alpha * np.dot(grad_q, grad_pi))
```

Nabin Katwal

10.13. Discount Factor (γ)

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}, \quad 0 \leq \gamma < 1$$

Explanation: The discount factor determines the weight given to future rewards. A smaller γ prioritizes immediate rewards, while a larger γ considers longer-term rewards.

Implementation:

```python
def discounted_return(rewards, gamma):
    G = 0
    for t, r in enumerate(rewards):
        G += (gamma**t) * r
    return G
```

Nabin Katwal

10.14. Expected SARSA

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[R_{t+1} + \gamma \mathbb{E}_{a'}[Q(s_{t+1}, a')] - Q(s_t, a_t) \right]$$

Explanation: Expected SARSA updates Q-values using the expected value of the next action, improving stability over standard SARSA.

Implementation:

```
def expected_sarsa(Q, state, action, reward, next_state, policy, alpha, gamma):
    expected_value = np.sum(policy[next_state] * Q[next_state])
    Q[state, action] += alpha * (reward + gamma * expected_value
    - Q[state, action])
```

Nabin Katwal

10.15.   Eligibility Traces Update (TD(λ))

$$\mathbf{e}_t = \gamma\lambda\mathbf{e}_{t-1} + \nabla_\theta V(s_t), \quad \theta \leftarrow \theta + \alpha\delta_t\mathbf{e}_t$$

Explanation: TD(λ) combines TD and Monte Carlo methods using eligibility traces, balancing bias and variance in value updates.

Implementation:

```
def td_lambda_update(V, eligibility, state, reward, next_state, alpha,
gamma, lambda_):
    delta = reward + gamma * V[next_state] - V[state]
    eligibility[state] += 1
    V += alpha * delta * eligibility
    eligibility *= gamma * lambda_
```

10.16.   TD Error

$$\delta_t = R_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

Explanation: The TD error measures the difference between predicted and observed rewards, guiding updates in temporal difference learning.

Implementation:

```
def td_error(V, state, reward, next_state, gamma):
    return reward + gamma * V[next_state] - V[state]
```

Nabin Katwal

10.17.   Stochastic Gradient Descent in RL

$$\theta \leftarrow \theta - \alpha \nabla_\theta \mathcal{L}(\theta)$$

Explanation: Stochastic gradient descent updates model parameters by minimizing a loss function, often used in function approximation for RL.

Implementation:

```
def sgd_update(theta, grad, alpha):
    return theta - alpha * grad
```

10.18.   Double Q-Learning

$$Q_1(s_t, a_t) \leftarrow Q_1(s_t, a_t) + \alpha \left[R_{t+1} + \gamma Q_2(s_{t+1}, \arg\max_a Q_1(s_{t+1}, a)) - Q_1(s_t, a_t) \right]$$

Explanation: Double Q-learning reduces overestimation bias by alternating updates between two Q-functions.

Implementation:

```
def double_q_learning_update(Q1, Q2, state, action, reward, next_state,
alpha, gamma):
    max_action = np.argmax(Q1[next_state])
    target = reward + gamma * Q2[next_state, max_action]
    Q1[state, action] += alpha * (target - Q1[state, action])
```

10.19.   Advantage Actor–Critic (A2C)

$$\delta_t = R_{t+1} + \gamma V(s_{t+1}) - V(s_t), \quad \theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(a_t \mid s_t)\delta_t$$

Explanation: A2C uses the advantage function to reduce variance in policy updates while learning the value function as a baseline.

Implementation:

```
def a2c_update(actor, critic, state, action, reward, next_state, alpha, gamma):
    delta = reward + gamma * critic[next_state] - critic[state]
    actor.update(state, action, alpha * delta)
    critic[state] += alpha * delta
```

10.20. Off-Policy Evaluation (Importance Sampling)

$$\mathbb{E}[\hat{G}] = \mathbb{E}\left[\prod_{t=0}^{T-1} \frac{\pi(a_t \mid s_t)}{\mu(a_t \mid s_t)} G_t\right]$$

Explanation: Importance sampling corrects for discrepancies between the behavior policy μ and the target policy π when estimating returns.

Implementation:

```
def importance_sampling(weights, returns):
    return np.sum(weights * returns)
```

10.21.   Policy Gradient Update Rule

$$\theta \leftarrow \theta + \alpha \nabla_\theta \mathbb{E}_{\pi_\theta} \left[G_t \log \pi_\theta(a_t \mid s_t) \right]$$

Explanation: The policy gradient algorithm updates parameters in the direction of performance improvement, directly optimizing the policy.

Implementation:

```
def policy_gradient_update(policy, rewards, states, actions, alpha):
    for state, action, reward in zip(states, actions, rewards):
        grad = policy.gradient(state, action)
        policy.update(state, action, alpha * reward * grad)
```

10.22. Soft Q-Learning Objective

$$\mathcal{L} = \mathbb{E}_{s,a} \left[Q(s,a) - \alpha \log \pi(a \mid s) \right]$$

Explanation: Soft Q-learning optimizes a policy by balancing reward maximization and entropy regularization.

Implementation:

```
def soft_q_update(Q, policy, state, action, reward, next_state, alpha, gamma):
    entropy = -policy.log_prob(action, state)
    target = reward + gamma * (Q[next_state].max() + alpha * entropy)
    Q[state, action] += alpha * (target - Q[state, action])
```

10.23.   Entropy-Regularized RL

$$\pi^* = \arg\max_{\pi} \mathbb{E}[G_t] + \alpha H(\pi)$$

Explanation: Entropy regularization encourages exploration by maximizing the entropy of the policy.

Implementation:

```
def entropy_regularized_update(policy, rewards, states,
actions, alpha, entropy_coeff):
    for state, action, reward in zip(states, actions, rewards):
        entropy = -policy.log_prob(action, state)
        grad = policy.gradient(state, action)
        policy.update(state, action, alpha *
        (reward + entropy_coeff * entropy) * grad)
```

Nabin Katwal

10.24.   Soft Actor–Critic (SAC)

$$\mathcal{L} = \mathbb{E}_{s,a}\left[Q(s,a) - \alpha \log \pi(a \mid s)\right], \quad Q(s,a) = R + \gamma V(s')$$

Explanation: SAC combines entropy regularization with actor–critic methods to improve stability and exploration in continuous control.

Implementation:

```
def sac_update(Q, policy, state, action, reward, next_state, alpha, gamma):
    entropy = -policy.log_prob(action, state)
    target = reward + gamma * (Q[next_state].max() + alpha * entropy)
    Q[state, action] += alpha * (target - Q[state, action])
```

10.25.   Trust Region Policy Optimization (TRPO)

$$\max_{\theta} \mathbb{E}_{\pi_\theta} \left[\frac{\pi_\theta(a \mid s)}{\pi_{\theta_{\text{old}}}(a \mid s)} A(s, a) \right], \quad \text{subject to } D_{\text{KL}}(\pi_\theta || \pi_{\theta_{\text{old}}}) \leq \delta$$